

Comprehensive Analysis & Delivery Time Estimation Strategy for Porter

Problem Statement:

Porter, India's largest intra-city logistics marketplace, needs to **predict accurate food delivery times** for orders placed through its platform.

The company partners with restaurants to deliver food directly to customers, and timely deliveries are critical for:

Customer satisfaction (avoiding frustration from delayed orders)

Operational efficiency (optimal delivery partner allocation)

Competitive advantage (transparent ETAs(Estimated Time of Arrivals) can differentiate Porter from rivals)

Understanding The Data

Start by importing necessary Python libraries like pandas, numpy, matplotlib, seaborn, etc. These help load the dataset, explore the structure, and perform analysis.

```
# importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# importing data and store in variable
PorterData = pd.read_csv("Porter_dataset.csv")

# check dataset dimensions (rows, columns)
f"Shape: {PorterData.shape}"

'Shape: (197428, 14)'
```

The dataset contains over 1.9 lakh + rows and 14 columns. It holds large-scale delivery/order information.

View the first few rows to understand how the data looks and what kind of values are stored.

```
# top 5 rows of data
PorterData.head()

  market_id      created_at actual_delivery_time \
0      1.0  2015-02-06 22:24:17  2015-02-06 23:27:16
1      2.0  2015-02-10 21:49:25  2015-02-10 22:56:29
2      3.0  2015-01-22 20:39:28  2015-01-22 21:09:09
```

```

3      3.0 2015-02-03 21:21:45 2015-02-03 22:13:00
4      3.0 2015-02-15 02:40:36 2015-02-15 03:20:26

          store_id store_primary_category
order_protocol \
0 df263d996281d984952c07998dc54358               american
1.0
1 f0ade77b43923b38237db569b016ba25             mexican
2.0
2 f0ade77b43923b38237db569b016ba25             NaN
1.0
3 f0ade77b43923b38237db569b016ba25             NaN
1.0
4 f0ade77b43923b38237db569b016ba25             NaN
1.0

      total_items  subtotal  num_distinct_items  min_item_price
max_item_price \
0           4        3441                  4            557
1239
1           1        1900                  1            1400
1400
2           1        1900                  1            1900
1900
3           6        6900                  5            600
1800
4           3        3900                  3            1100
1600

      total_onshift_partners  total_busy_partners
total_outstanding_orders
0                   33.0              14.0
21.0
1                   1.0               2.0
2.0
2                   1.0               0.0
0.0
3                   1.0               1.0
2.0
4                   6.0               6.0
9.0

```

Column Description,

1. market_id: Integer ID for the market where the restaurant is located
2. created_at: Timestamp at which the order was placed
3. actual_delivery_time: Timestamp when the order was delivered
4. store_primary_category: Category of the restaurant

5. order_protocol: Integer code value for the order protocol (e.g., through Porter, call to restaurant, pre-booked, third-party, etc.)
6. total_items: Total number of items in the order
7. subtotal: Final price of the order
8. num_distinct_items: Number of distinct items in the order
9. min_item_price: Price of the cheapest item in the order
10. max_item_price: Price of the most expensive item in the order
11. total_onshift_partners: Number of delivery partners on duty when the order was placed
12. total_busy_partners: Number of delivery partners attending to other tasks
13. total_outstanding_orders: Total number of orders to be fulfilled at that moment

Checking Column Variables and Types.

```
# printing information about the data
PorterData.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   market_id        196441 non-null   float64 
 1   created_at       197428 non-null   object  
 2   actual_delivery_time 197421 non-null   object  
 3   store_id         197428 non-null   object  
 4   store_primary_category 192668 non-null   object  
 5   order_protocol    196433 non-null   float64 
 6   total_items       197428 non-null   int64   
 7   subtotal          197428 non-null   int64   
 8   num_distinct_items 197428 non-null   int64   
 9   min_item_price    197428 non-null   int64   
 10  max_item_price    197428 non-null   int64   
 11  total_onshift_partners 181166 non-null   float64 
 12  total_busy_partners 181166 non-null   float64 
 13  total_outstanding_orders 181166 non-null   float64 
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB
```

Look at the names of all columns and their current data types (like object, int, float). Some columns like created_at need to be converted to date format for better analysis.

Check for missing values in each column to understand where data is incomplete.

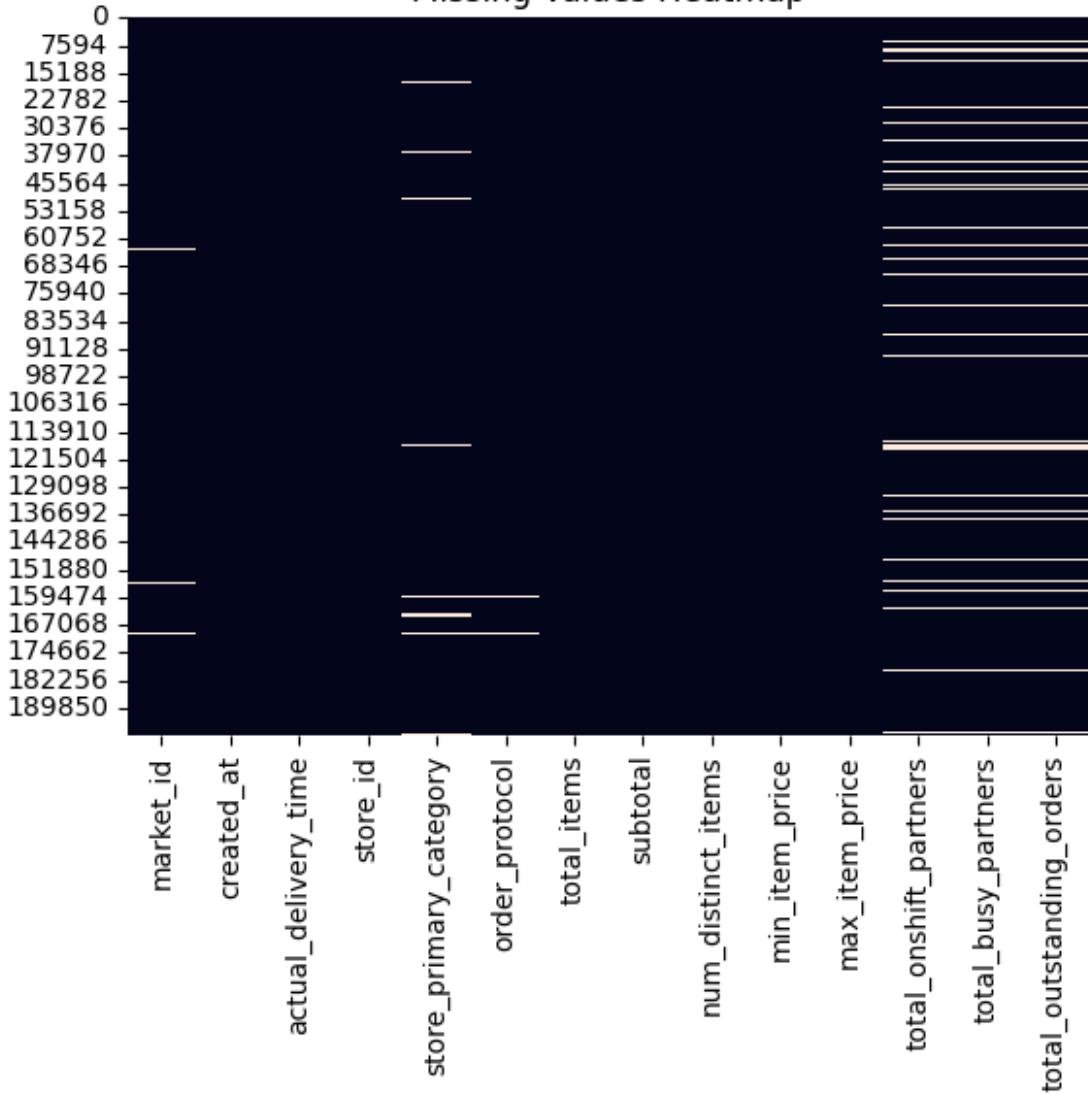
```
# total missing values per column
PorterData.isnull().sum()

market_id            987
created_at           0
actual_delivery_time 7
```

```
store_id                      0
store_primary_category      4760
order_protocol                  995
total_items                      0
subtotal                          0
num_distinct_items                 0
min_item_price                     0
max_item_price                     0
total_onshift_partners        16262
total_busy_partners            16262
total_outstanding_orders       16262
dtype: int64

# visualize missing data
sns.heatmap(PorterData.isnull(), cbar=False)
plt.title("Missing Values Heatmap")
plt.show()
```

Missing Values Heatmap



```
# check null percentage
PorterData.isnull().mean() * 100
```

market_id	0.499929
created_at	0.000000
actual_delivery_time	0.003546
store_id	0.000000
store_primary_category	2.411006
order_protocol	0.503981
total_items	0.000000
subtotal	0.000000
num_distinct_items	0.000000
min_item_price	0.000000
max_item_price	0.000000
total_onshift_partners	8.236927

```

total_busy_partners      8.236927
total_outstanding_orders 8.236927
dtype: float64

```

Columns like total_onshift_partners, total_busy_partners, and total_outstanding_orders have a large number of nulls.

Other columns like market_id, actual_delivery_time, store_delivery_time, store_primary_category, and order_protocol have fewer null values.

Exploratory Data Analysis (EDA)

View summary statistics (mean, median, max, min, etc.) for each column to get a quick understanding of the dataset

```

# generating descriptive statistics
PorterData.describe()

      market_id  order_protocol  total_items  subtotal \
count  196441.000000  196433.000000  197428.000000  197428.000000
mean   2.978706     2.882352     3.196391    2682.331402
std    1.524867     1.503771     2.666546    1823.093688
min    1.000000     1.000000     1.000000     0.000000
25%   2.000000     1.000000     2.000000    1400.000000
50%   3.000000     3.000000     3.000000    2200.000000
75%   4.000000     4.000000     4.000000    3395.000000
max   6.000000     7.000000    411.000000   27100.000000

      num_distinct_items  min_item_price  max_item_price \
count  197428.000000  197428.000000  197428.000000
mean   2.670791       686.218470    1159.588630
std    1.630255       522.038648    558.411377
min    1.000000      -86.000000     0.000000
25%   1.000000      299.000000    800.000000
50%   2.000000      595.000000   1095.000000
75%   3.000000      949.000000   1395.000000
max   20.000000     14700.000000  14700.000000

      total_onshift_partners  total_busy_partners
total_outstanding_orders
count          181166.000000           181166.000000
181166.000000
mean            44.808093            41.739747
58.050065
std             34.526783            32.145733
52.661830
min            -4.000000            -5.000000
6.000000
25%            17.000000            15.000000
17.000000

```

50%	37.000000	34.000000
41.000000		
75%	65.000000	62.000000
85.000000		
max	171.000000	154.000000
285.000000		

Total Items Column:

- 75% of orders have 4 or fewer items.
- Maximum items go above 400 – likely outliers.

Item Price Columns (Min & Max):

- Both have unusually high values (up to 14,700), which seems unrealistic for a single item.
- Some values are negative, suggesting data entry mistakes.

Subtotal Column:

- Some subtotal values are zero or unusually high (up to 27,100).
- Could be due to missing transaction values or calculation errors.

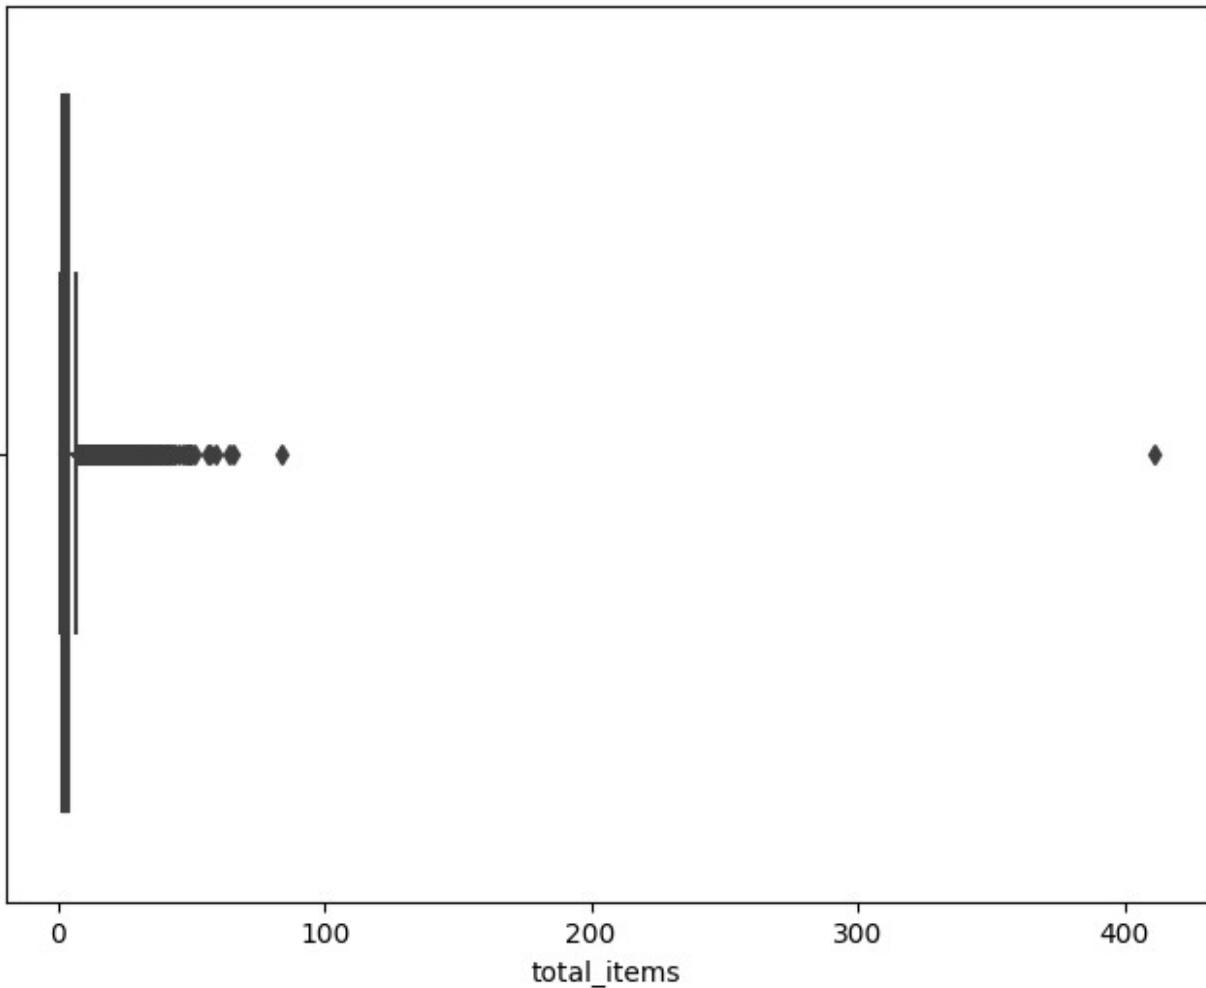
Logical Inconsistencies:

- In a few rows, total_busy_partners is greater than total_onshift_partners, which shouldn't happen.
- In some rows, all three columns (onshift, busy, outstanding_orders) are zero, but still, there's a delivery recorded. That's not possible.

Outlier & Pattern Detection

```
# boxplot for total_items column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['total_items'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



```
PorterData[(PorterData['total_items'] > 50)]
```

```
market_id      created_at actual_delivery_time \
11055        2.0 2015-01-22 01:07:03 2015-01-22 01:55:43
15053        4.0 2015-01-26 22:48:56 2015-01-26 23:28:02
47231        2.0 2015-02-06 00:42:39 2015-02-06 01:33:34
75577        1.0 2015-01-23 04:33:00 2015-01-23 05:07:45
105348       3.0 2015-01-23 15:46:35 2015-01-23 16:56:32
182223       6.0 2015-02-15 19:39:32 2015-02-15 20:54:10
182796       2.0 2015-02-17 05:45:05 2015-02-17 06:14:01
182800       2.0 2015-02-18 05:42:03 2015-02-18 06:11:25

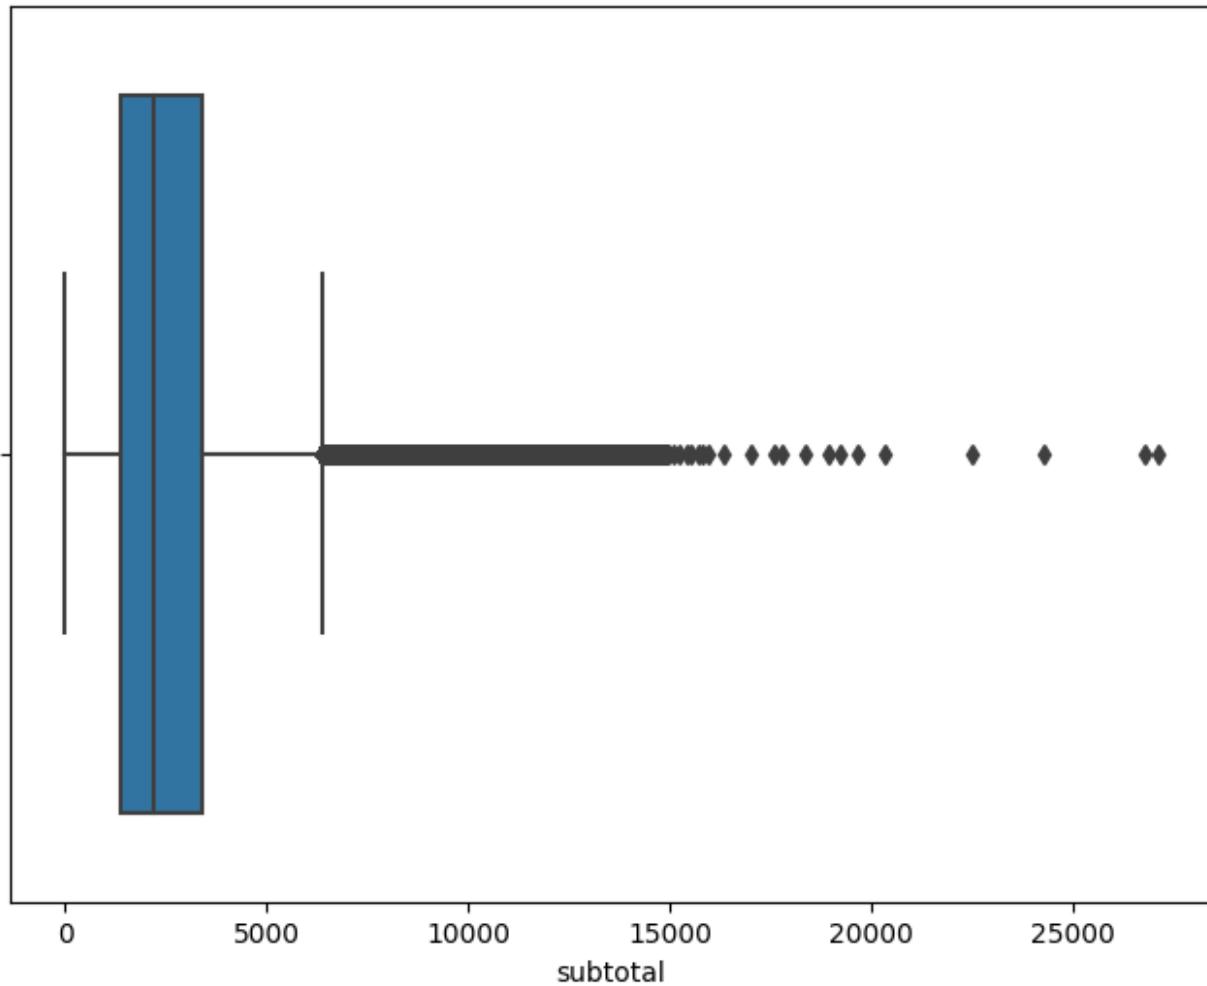
                           store_id store_primary_category \
11055  c8512d142a2d849725f31a9a7a361ab9                  fast
15053  4a308d84cdd04aa2015bbe13622d5d7c                  fast
47231  c156cea027720c227089e679b3ae9d1b                  fast
75577  452e91de642a8e9c43121664d5d3c05c                  fast
105348 d14c2267d848abeb81fd590f371d39bd                  fast
```

182223	e50372d3fee4eadec9c42aa6528097cc	fast			
182796	42a3ddf2e1df611a280d556f1c81996a	fast			
182800	42a3ddf2e1df611a280d556f1c81996a	fast			
<hr/>					
11055	order_protocol	total_items	subtotal	num_distinct_items	\
11055	4.0	56	2317		7
15053	4.0	57	858		5
47231	4.0	411	3115		5
75577	4.0	59	2911		3
105348	4.0	51	1343		4
182223	4.0	84	1016		4
182796	4.0	64	1166		4
182800	4.0	66	1634		6
<hr/>					
11055	min_item_price	max_item_price	total_onshift_partners	\	
11055	0	289		35.0	
15053	0	229		25.0	
47231	0	299		35.0	
75577	0	329		25.0	
105348	0	169		3.0	
182223	0	289		NaN	
182796	0	499		28.0	
182800	0	499		30.0	
<hr/>					
11055	total_busy_partners	total_outstanding_orders			
11055	27.0	26.0			
15053	54.0	19.0			
47231	35.0	39.0			
75577	19.0	20.0			
105348	1.0	2.0			
182223	NaN	NaN			
182796	19.0	21.0			
182800	28.0	32.0			

Orders with more than 50 items usually come from stores with 'fast' primary category and order_protocol = 4.

```
# boxplot for subtotal column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['subtotal'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



```
PorterData[(PorterData['subtotal'] > 15000)]
```

	market_id	created_at	actual_delivery_time	\
8994	1.0	2015-01-27 20:08:26	2015-01-27 21:21:37	
19966	4.0	2015-02-16 01:12:30	2015-02-16 02:19:05	
37806	4.0	2015-02-13 21:20:15	2015-02-13 22:23:24	
44705	4.0	2015-01-29 21:02:27	2015-01-29 21:42:20	
53072	1.0	2015-02-05 01:02:41	2015-02-05 02:25:20	
54486	4.0	2015-01-22 23:59:18	2015-01-23 00:39:52	
69243	1.0	2015-01-29 02:33:00	2015-01-29 04:03:25	
71316	6.0	2015-02-12 00:50:03	2015-02-12 02:07:01	
71702	6.0	2015-01-27 02:22:26	2015-01-27 03:47:09	
73129	4.0	2015-02-10 01:42:46	2015-02-10 02:36:05	
82763	6.0	2015-02-12 01:28:39	2015-02-12 03:39:57	
83330	6.0	2015-01-28 02:12:01	2015-01-28 03:50:48	
96456	1.0	2015-02-04 19:21:43	2015-02-04 20:20:25	
140458	4.0	2015-02-18 01:16:47	2015-02-18 02:29:14	
152185	4.0	2015-01-22 00:09:59	2015-01-22 01:29:21	

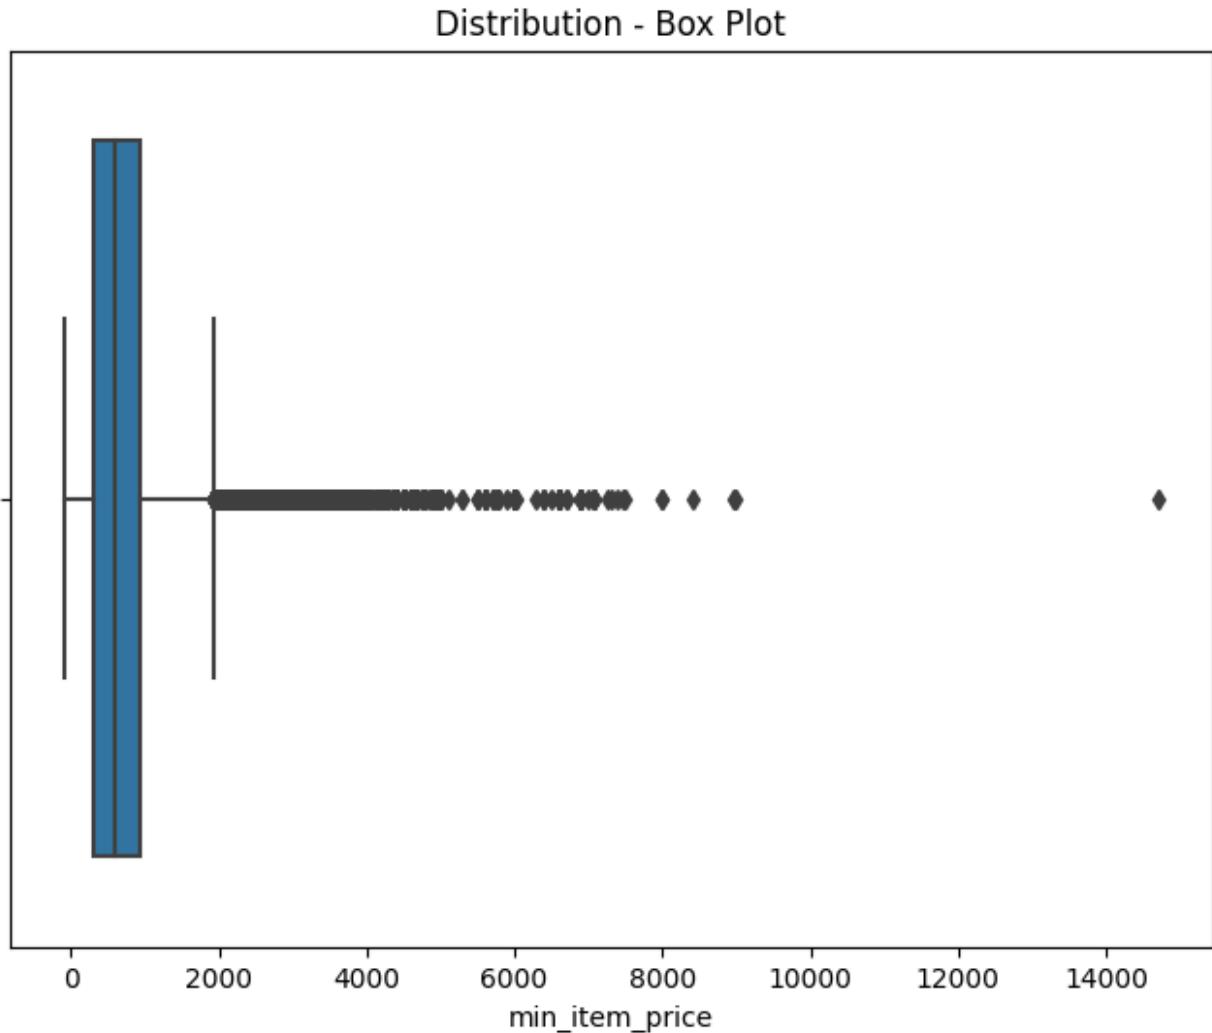
171276	4.0	2015-02-02	17:46:21	2015-02-02	18:46:36	
173912	6.0	2015-02-13	16:29:21	2015-02-13	18:27:58	
176874	4.0	2015-02-13	20:44:02	2015-02-13	22:00:09	
181734	2.0	2015-02-16	02:22:47	2015-02-16	03:42:16	
188378	1.0	2015-01-27	02:11:51	2015-01-27	03:35:26	
				store_id	store_primary_category	\
8994	652cf38361a209088302ba2b8b7f51e0				salad	
19966	03afdbd66e7929b125f8597834fa83a4				sushi	
37806	b9a25e422ba96f7572089a00b838c3f8				indian	
44705	09fb05dd477d4ae6479985ca56c5a12d				italian	
53072	28b60a16b55fd531047c0c958ce14b95				pizza	
54486	ccc0aa1b81bf81e16c676ddb977c5881				seafood	
69243	7f6b550688de8db0290009ba6abd673c				NaN	
71316	9cea886b9f44a3c2df1163730ab64994				greek	
71702	c900fe92840c527a0c54f28640c2f254				catering	
73129	dbe40cccd11b1fb869099e58e00076027				vietnamese	
82763	f58c9875ac84dfe1fbe91b918773d050				japanese	
83330	f58c9875ac84dfe1fbe91b918773d050				japanese	
96456	748ba69d3e8d1af87f84fee909eef339				italian	
140458	e2ef524fbf3d9fe611d5a8e90fefdc9c				american	
152185	9cea10c7ff109c6e61727a0d45492ead				vietnamese	
171276	3dd48ab31d016ffcbf3314df2b3cb9ce				greek	
173912	4607f7fff0dce694258e1c637512aa9d				sandwich	
176874	f1b6f2857fb6d44dd73c7041e0aa0f19				middle-eastern	
181734	91665c93b72f55b2e4f1048fc8289d04				smoothie	
188378	fd2c5e4680d9a01dba3aada5ece22270				vietnamese	
	order_protocol	total_items	subtotal	num_distinct_items	\	
8994	1.0	21	18370	11		
19966	3.0	19	15505	12		
37806	3.0	28	15238	9		
44705	3.0	10	15089	8		
53072	5.0	20	20350	14		
54486	5.0	14	16350	11		
69243	4.0	24	17008	13		
71316	5.0	11	15800	9		
71702	3.0	18	24300	14		
73129	1.0	11	19250	11		
82763	1.0	27	27100	19		
83330	1.0	20	19650	12		
96456	5.0	17	22500	14		
140458	3.0	6	15435	6		
152185	5.0	18	15710	14		
171276	3.0	16	18920	4		
173912	4.0	17	17810	14		
176874	5.0	14	17600	13		
181734	2.0	4	15960	2		
188378	5.0	25	26800	13		

	min_item_price	max_item_price	total_onshift_partners	\
8994	350	1380	13.0	
19966	310	2400	84.0	
37806	245	1399	38.0	
44705	1130	1710	80.0	
53072	450	1250	7.0	
54486	300	1650	37.0	
69243	61	1556	55.0	
71316	350	2000	16.0	
71702	600	2200	NaN	
73129	200	3100	75.0	
82763	200	3000	NaN	
83330	200	2000	NaN	
96456	450	1750	18.0	
140458	795	3570	76.0	
152185	265	1275	33.0	
171276	695	1695	10.0	
173912	375	1295	NaN	
176874	400	1700	63.0	
181734	3990	3990	56.0	
188378	200	1800	13.0	
	total_busy_partners	total_outstanding_orders		
8994	13.0	12.0		
19966	80.0	142.0		
37806	20.0	22.0		
44705	64.0	70.0		
53072	8.0	11.0		
54486	18.0	19.0		
69243	49.0	63.0		
71316	6.0	5.0		
71702	NaN	NaN		
73129	67.0	92.0		
82763	NaN	NaN		
83330	NaN	NaN		
96456	17.0	21.0		
140458	75.0	105.0		
152185	14.0	14.0		
171276	11.0	14.0		
173912	NaN	NaN		
176874	63.0	89.0		
181734	53.0	92.0		
188378	14.0	17.0		

Oddly, even when items are more, subtotal is sometimes less, which looks unrealistic.

```
# boxplot for min_item_price column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['min_item_price'])
```

```
plt.title('Distribution - Box Plot')
plt.show()
```



```
PorterData[(PorterData['min_item_price'] > 5000)]
```

	market_id	created_at	actual_delivery_time	\
999	2.0	2015-01-28 00:45:38	2015-01-28 01:15:26	
6814	2.0	2015-01-24 05:10:20	2015-01-24 06:05:35	
7545	6.0	2015-02-13 18:34:39	2015-02-13 19:36:52	
7899	6.0	2015-02-07 03:52:47	2015-02-07 05:11:13	
13548	4.0	2015-01-30 22:16:08	2015-01-30 23:06:16	
...
184000	2.0	2015-02-03 01:23:48	2015-02-03 02:18:14	
188815	5.0	2015-01-22 02:27:31	2015-01-22 03:08:34	
190087	2.0	2015-02-01 03:46:30	2015-02-01 04:33:39	
190213	2.0	2015-01-29 02:56:02	2015-01-29 03:39:58	
195513	2.0	2015-02-14 03:38:56	2015-02-14 04:40:42	

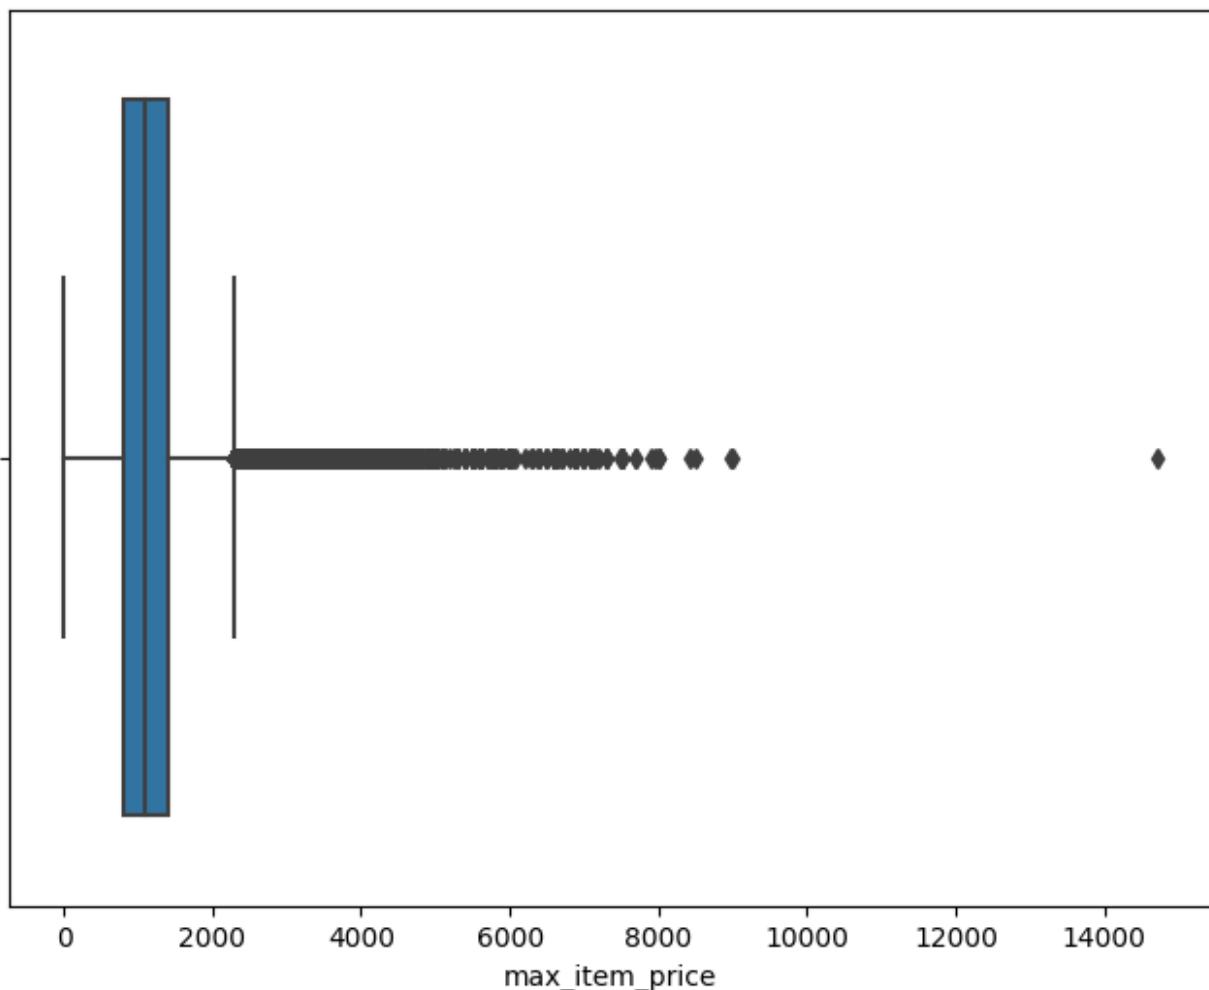
	store_id	store_primary_category	\	
999	0bb0846327772451045bd30dd347821b	barbecue		
6814	140f6969d5213fd0ece03148e62e461e	japanese		
7545	f3a493fb0b6dc6357b9d89d6bdc1f2af	dessert		
7899	8cbd005a556ccd4211ce43f309bc0eac	thai		
13548	e723e2ae3e04e8028e119ee592e81974	NaN		
...	
184000	777125b977ddc0317d0533782d3c27b5	chinese		
188815	15b33319db65d343906d085ba0500783	dessert		
190087	9f319422ca17b1082ea49820353f14ab	american		
190213	9f319422ca17b1082ea49820353f14ab	american		
195513	b783acd4479bf1b8a981bb023b363043	american		
	order_protocol	total_items	subtotal	
		num_distinct_items	\	
999	3.0	1	7475	1
6814	2.0	1	6400	1
7545	2.0	1	7080	1
7899	3.0	1	9900	1
13548	1.0	1	7299	1
...
184000	3.0	1	6500	1
188815	2.0	1	6000	1
190087	3.0	1	7029	1
190213	3.0	1	6889	1
195513	1.0	1	8959	1
	min_item_price	max_item_price	total_onshift_partners	
			\	
999	7475	7475	15.0	
6814	6400	6400	59.0	
7545	7080	7080	NaN	
7899	6000	6000	NaN	
13548	7299	7299	7.0	
...	
184000	6500	6500	49.0	
188815	6000	6000	18.0	
190087	6889	6889	101.0	
190213	6889	6889	103.0	
195513	8959	8959	143.0	
	total_busy_partners	total_outstanding_orders		
			\	
999	14.0	14.0		
6814	69.0	111.0		
7545	NaN	NaN		
7899	NaN	NaN		
13548	7.0	13.0		
...		
184000	31.0	55.0		
188815	13.0	20.0		
190087	102.0	184.0		

```
190213           87.0          127.0
195513          131.0          274.0
```

```
[72 rows x 14 columns]
```

```
# boxplot for max_item_price column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['max_item_price'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



```
PorterData[(PorterData['max_item_price'] > 6000)]
```

```
   market_id      created_at actual_delivery_time \
999      2.0 2015-01-28 00:45:38 2015-01-28 01:15:26
6814     2.0 2015-01-24 05:10:20 2015-01-24 06:05:35
7545     6.0 2015-02-13 18:34:39 2015-02-13 19:36:52
```

12790	2.0	2015-02-14	20:17:26	2015-02-14	21:06:35
13548	4.0	2015-01-30	22:16:08	2015-01-30	23:06:16
...
182045	3.0	2015-01-26	01:21:21	2015-01-26	02:38:45
184000	2.0	2015-02-03	01:23:48	2015-02-03	02:18:14
190087	2.0	2015-02-01	03:46:30	2015-02-01	04:33:39
190213	2.0	2015-01-29	02:56:02	2015-01-29	03:39:58
195513	2.0	2015-02-14	03:38:56	2015-02-14	04:40:42
store_id store_primary_category \					
999	0bb0846327772451045bd30dd347821b				barbecue
6814	140f6969d5213fd0ece03148e62e461e				japanese
7545	f3a493fb0b6dc6357b9d89d6bdc1f2af				dessert
12790	9854d7afce413aa13cd0a1d39d0bcec5				american
13548	e723e2ae3e04e8028e119ee592e81974				NaN
...
182045	38181d991caac98be8fb2ecb8bd0f166				american
184000	777125b977ddc0317d0533782d3c27b5				chinese
190087	9f319422ca17b1082ea49820353f14ab				american
190213	9f319422ca17b1082ea49820353f14ab				american
195513	b783acd4479bf1b8a981bb023b363043				american
order_protocol total_items subtotal num_distinct_items \					
999	3.0	1	7475		1
6814	2.0	1	6400		1
7545	2.0	1	7080		1
12790	1.0	2	8048		2
13548	1.0	1	7299		1
...
182045	4.0	1	6699		1
184000	3.0	1	6500		1
190087	3.0	1	7029		1
190213	3.0	1	6889		1
195513	1.0	1	8959		1
min_item_price max_item_price total_onshift_partners \					
999	7475	7475		15.0	
6814	6400	6400		59.0	
7545	7080	7080		NaN	
12790	79	7699		70.0	
13548	7299	7299		7.0	
...
182045	6699	6699		41.0	
184000	6500	6500		49.0	
190087	6889	6889		101.0	
190213	6889	6889		103.0	
195513	8959	8959		143.0	
total_busy_partners total_outstanding_orders					
999		14.0		14.0	

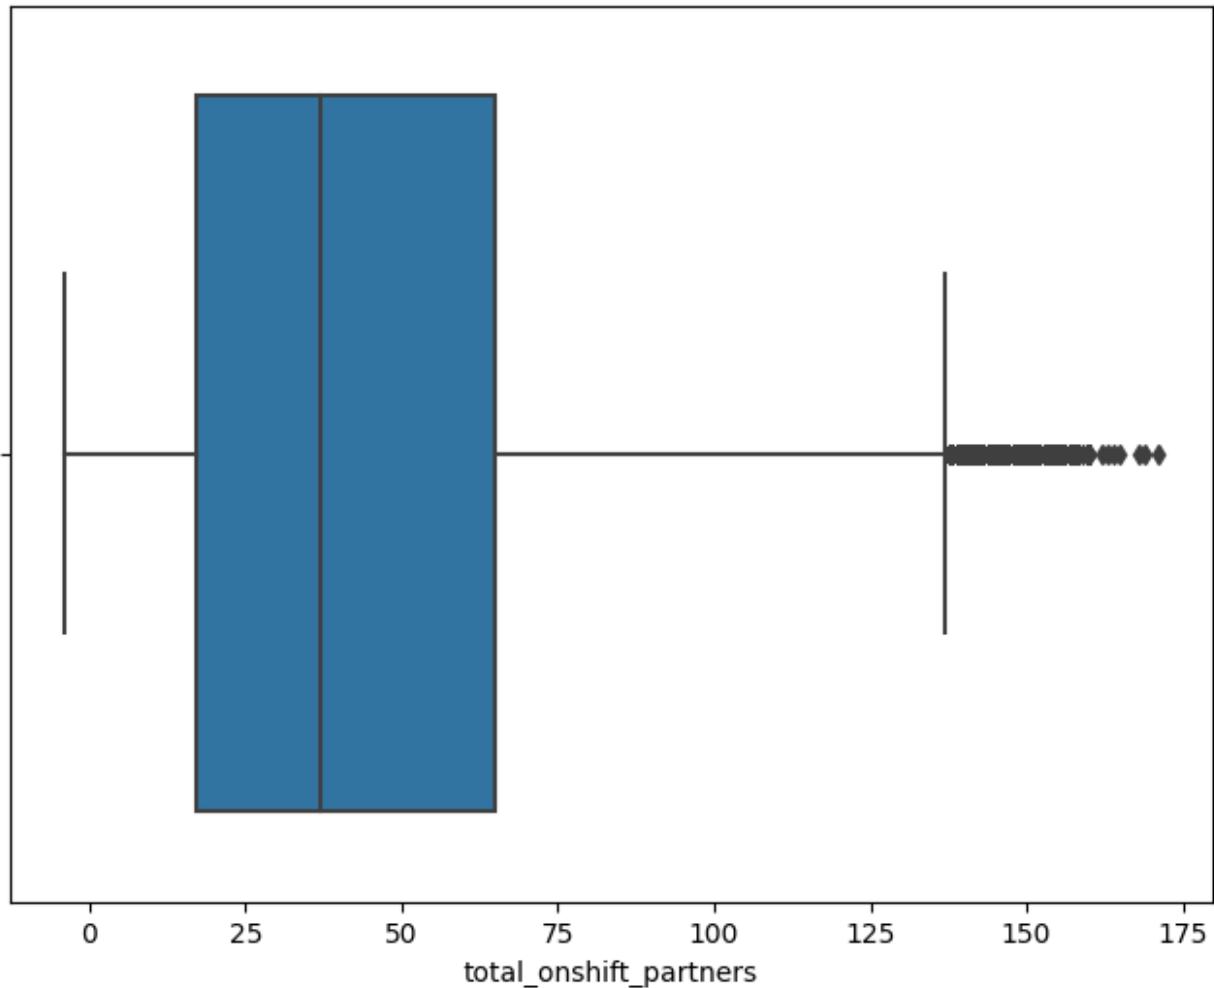
6814	69.0	111.0
7545	NaN	NaN
12790	66.0	71.0
13548	7.0	13.0
...
182045	40.0	47.0
184000	31.0	55.0
190087	102.0	184.0
190213	87.0	127.0
195513	131.0	274.0

[71 rows x 14 columns]

High max and min item prices (14,700) are seen for the breakfast store category, often with only one item in the order.

```
# boxplot for total_onshift_partners column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['total_onshift_partners'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



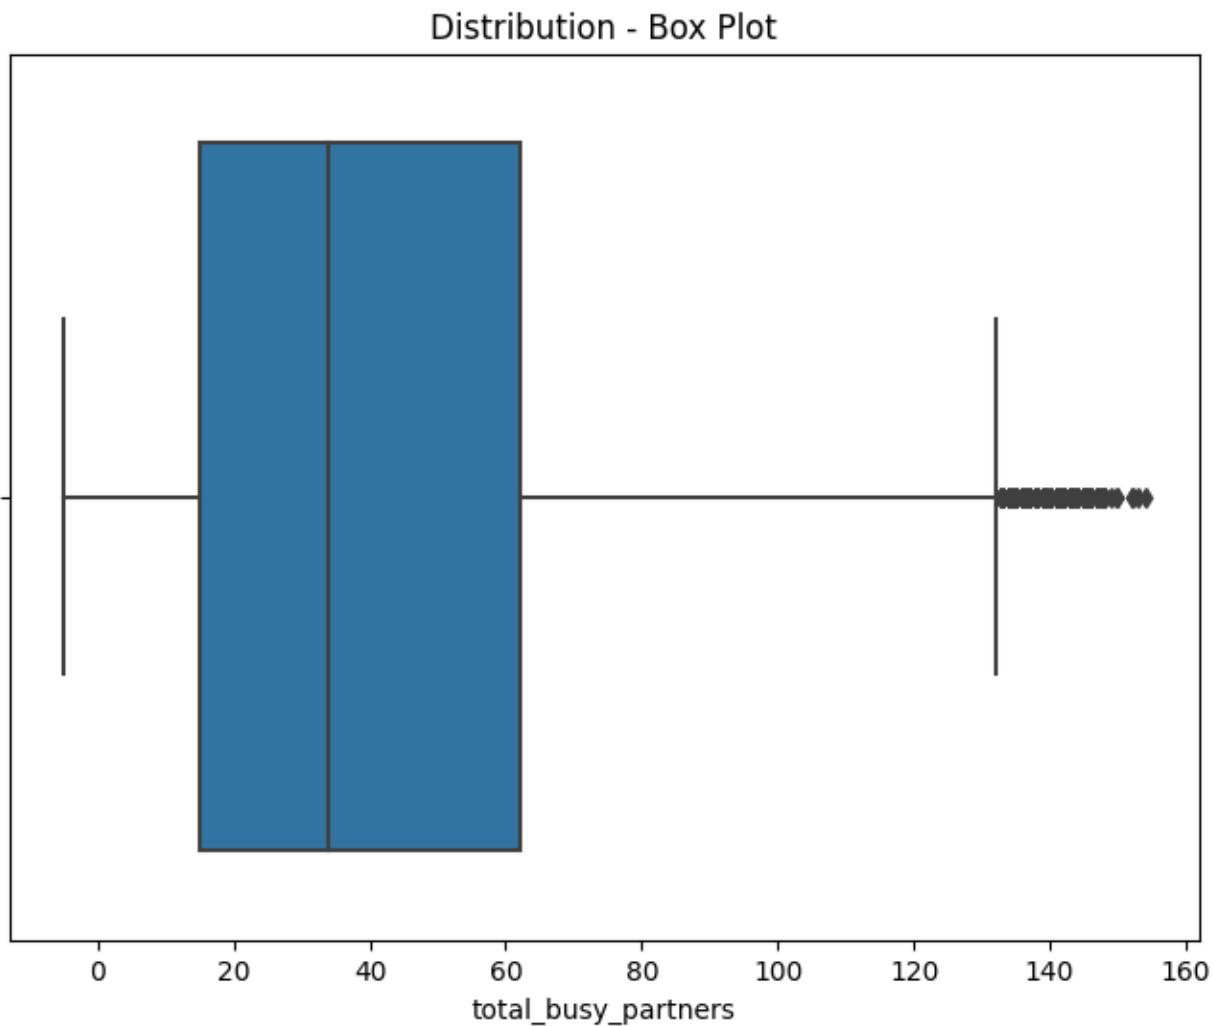
```
PorterData[(PorterData['total_onshift_partners'] > 150)]
```

	market_id	created_at	actual_delivery_time	\
1679	2.0	2015-02-07 02:45:20	2015-02-07 03:25:44	
3328	2.0	2015-02-06 02:27:12	2015-02-06 03:00:02	
4101	2.0	2015-02-14 02:56:43	2015-02-14 04:08:33	
5780	2.0	2015-02-07 02:56:37	2015-02-07 03:38:35	
5930	2.0	2015-02-07 02:46:03	2015-02-07 03:27:42	
...
194419	2.0	2015-02-07 02:44:51	2015-02-07 04:05:32	
195448	2.0	2015-02-06 02:42:42	2015-02-06 03:27:28	
195971	2.0	2015-02-07 02:57:05	2015-02-07 03:36:47	
196128	2.0	2015-02-07 02:31:28	2015-02-07 03:15:44	
196180	2.0	2015-02-07 02:58:55	2015-02-07 03:42:06	
		store_id	store_primary_category	\
1679	70821a40b06f8751781d5a895357da67		chinese	
3328	150784e5fbef562400a0cd1111471d6a		american	

4101	150784e5fbef562400a0cd1111471d6a	american
5780	57c0531e13f40b91b3b0f1a30b529a1d	pizza
5930	115f89503138416a242f40fb7d7f338e	pasta
...
194419	371bce7dc83817b7893bcdeed13799b5	pizza
195448	b783acd4479bf1b8a981bb023b363043	american
195971	84d9ee44e457ddef7f2c4f25dc8fa865	sandwich
196128	84d9ee44e457ddef7f2c4f25dc8fa865	sandwich
196180	84d9ee44e457ddef7f2c4f25dc8fa865	sandwich
order_protocol \		
1679	1.0	4
3328	1.0	1
4101	1.0	1
5780	5.0	1
5930	3.0	4
...
194419	3.0	4
195448	1.0	1
195971	5.0	4
196128	5.0	2
196180	5.0	4
min_item_price \		
1679	595	1275
3328	949	949
4101	1899	1899
5780	1645	1645
5930	250	925
...
194419	995	1295
195448	1619	1619
195971	100	999
196128	999	999
196180	295	1111
total_onshift_partners \		
1679	152.0	152.0
3328	155.0	155.0
4101	151.0	151.0
5780	156.0	156.0
5930	152.0	152.0
...
194419	154.0	154.0
195448	157.0	157.0
195971	156.0	156.0
196128	151.0	151.0
196180	156.0	156.0
total_busy_partners \ total_outstanding_orders		
1679	131.0	211.0
3328	76.0	114.0
4101	127.0	231.0
5780	137.0	215.0
5930	131.0	211.0
...
194419	129.0	199.0
195448	87.0	138.0
195971	137.0	215.0
196128	114.0	195.0
196180	137.0	215.0

[330 rows x 14 columns]

```
# boxplot for total_busy_partners column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['total_busy_partners'])
plt.title('Distribution - Box Plot')
plt.show()
```



```
PorterData[(PorterData['total_busy_partners'] > 140)]
   market_id      created_at actual_delivery_time \
1266      4.0 2015-01-25 02:26:42 2015-01-25 03:07:42
1464      4.0 2015-01-25 02:14:50 2015-01-25 02:51:20
1843      4.0 2015-01-25 03:00:26 2015-01-25 04:14:02
1853      4.0 2015-01-25 02:39:24 2015-01-25 03:36:16
1905      4.0 2015-01-25 02:38:29 2015-01-25 03:39:37
...
196453     2.0 2015-02-14 03:56:01 2015-02-14 04:48:35
197097     4.0 2015-01-25 03:14:04 2015-01-25 04:10:35
197104     4.0 2015-01-25 02:20:30 2015-01-25 02:56:05
```

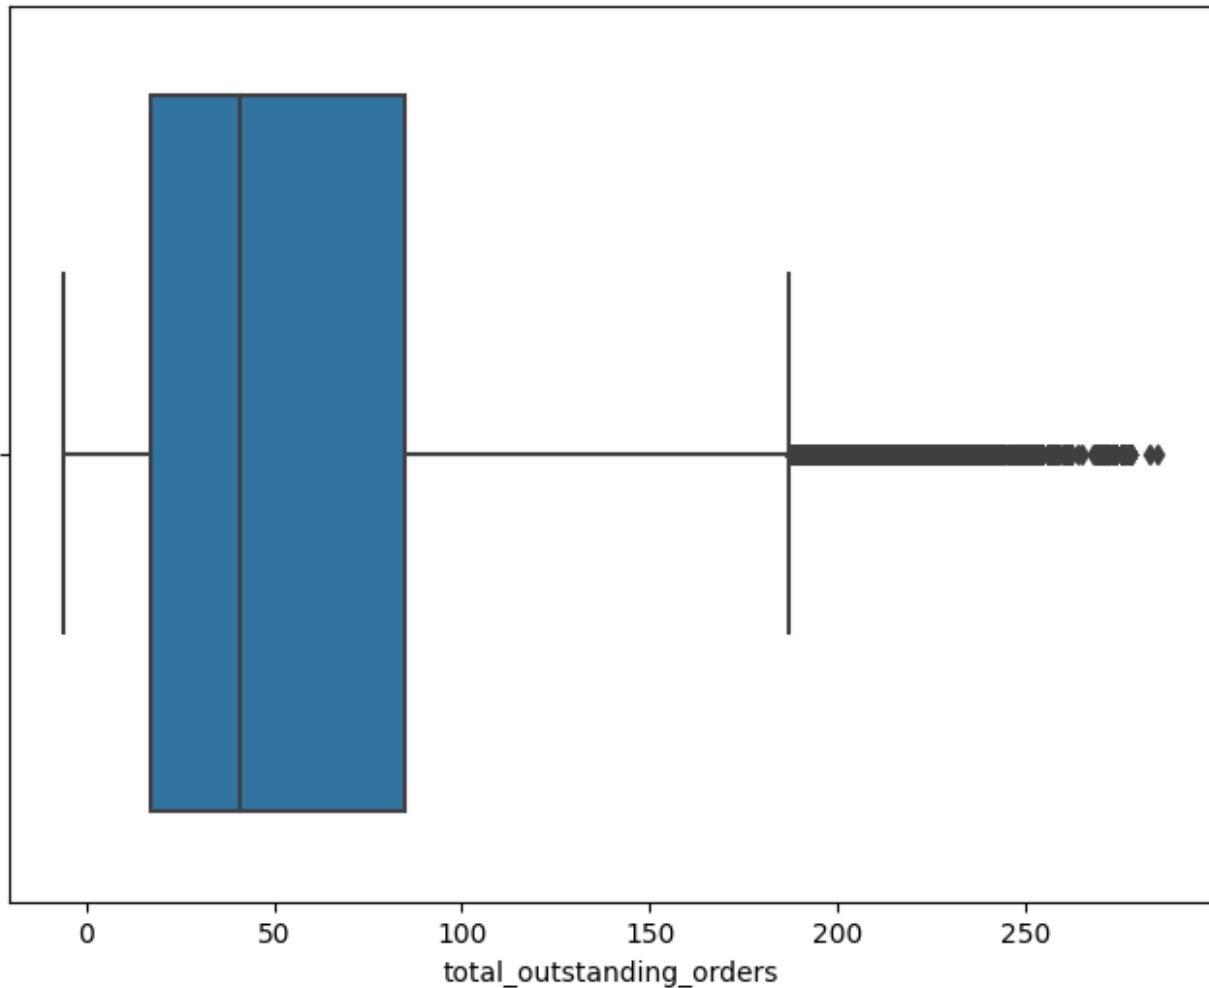
197328	4.0	2015-01-25 02:42:49	2015-01-25 03:29:09
197338	4.0	2015-01-25 03:13:33	2015-01-25 04:24:24
1266	f7177163c833dff4b38fc8d2872f1ec6	store_id	store_primary_category \
1464	6ea9ab1baa0efb9e19094440c317e21b		dessert
1843	3f088ebeda03513be71d34d214291986		italian
1853	3f088ebeda03513be71d34d214291986		mexican
1905	6ea9ab1baa0efb9e19094440c317e21b		mexican
...
196453	0d1a9651497a38d8b1c3871c84528bd4		italian
197097	17e62166fc8586dfa4d1bc0e1742c08b		other
197104	17e62166fc8586dfa4d1bc0e1742c08b		vietnamese
197328	4b04a686b0ad13dce35fa99fa4161c65		vietnamese
197338	4b04a686b0ad13dce35fa99fa4161c65		italian
1266	5.0	order_protocol	total_items \
1464	3.0		subtotal
1843	2.0		num_distinct_items
1853	2.0		4
1905	3.0		2
...	...		1
196453	3.0		1
197097	5.0		1400
197104	5.0		1400
197328	2.0		1
197338	2.0		1
1266	5.0	1905	2
1464	3.0	196453	2
1843	2.0	197097	2
1853	2.0	197104	1
1905	3.0	197328	1
...	...	197338	1
196453	3.0	2.0	2
197097	5.0	2.0	2
197104	5.0	2.0	2
197328	2.0	2.0	2
197338	2.0	2.0	2
1266	795	min_item_price	total_onshift_partners \
1464	1575		103.0
1843	1400		103.0
1853	1400		108.0
1905	575		106.0
...	...		106.0
196453	305		130.0
197097	750		103.0
197104	795		105.0
197328	1148		106.0
197338	1825		106.0
1266	795	196453	103.0
1464	1575	197097	103.0
1843	1400	197104	105.0
1853	1400	197328	106.0
1905	575	197338	106.0
...	...	1266	103.0
196453	305	1464	103.0
197097	750	1843	105.0
197104	795	1853	106.0
197328	1148	1905	106.0
197338	1825	196453	103.0
1266	143.0	total_busy_partners	total_outstanding_orders
1464	146.0		136.0
1843	148.0		134.0
1853	146.0		124.0
1905	146.0		138.0
...	...		138.0
196453	141.0		260.0

```
197097           147.0            120.0
197104           142.0            122.0
197328           144.0            137.0
197338           147.0            120.0
```

[219 rows x 14 columns]

```
# boxplot for total_outstanding_orders column
plt.figure(figsize=(8, 6))
sns.boxplot(x=PorterData['total_outstanding_orders'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



```
PorterData[(PorterData['total_outstanding_orders'] > 250)]
   market_id      created_at actual_delivery_time \
170        4.0  2015-02-07 02:34:41  2015-02-07 03:20:31
```

1680	2.0	2015-02-14 03:26:43	2015-02-14 04:08:55
2319	2.0	2015-02-14 03:44:27	2015-02-14 04:51:51
3000	4.0	2015-02-07 02:39:13	2015-02-07 03:51:31
3004	4.0	2015-02-07 02:31:31	2015-02-07 04:03:10
...
195905	2.0	2015-01-31 03:05:11	2015-01-31 04:14:18
196141	2.0	2015-01-24 03:09:38	2015-01-24 04:27:04
196453	2.0	2015-02-14 03:56:01	2015-02-14 04:48:35
196469	2.0	2015-02-14 03:18:25	2015-02-14 04:01:40
196506	2.0	2015-02-14 03:54:01	2015-02-14 05:14:27
store_id store_primary_category \			
170	a87ff679a2f3e71d9181a67b7542122c		mediterranean
1680	70821a40b06f8751781d5a895357da67		chinese
2319	f806c5d2707545d718717be03e69a8d4		fast
3000	d0cbf1ala1aa1726784df15a81ead214f7		italian
3004	d0cbf1ala1aa1726784df15a81ead214f7		greek
...
195905	84d9ee44e457ddef7f2c4f25dc8fa865		sandwich
196141	84d9ee44e457ddef7f2c4f25dc8fa865		sandwich
196453	0d1a9651497a38d8b1c3871c84528bd4		other
196469	0d1a9651497a38d8b1c3871c84528bd4		other
196506	0d1a9651497a38d8b1c3871c84528bd4		other
order_protocol total_items subtotal num_distinct_items \			
170	3.0	4	2945
1680	1.0	2	1545
2319	4.0	6	1555
3000	5.0	4	4396
3004	5.0	3	4396
...
195905	3.0	3	3219
196141	5.0	2	2290
196453	3.0	2	1549
196469	3.0	3	1947
196506	3.0	7	3193
min_item_price max_item_price total_onshift_partners \			
170	425	1245	129.0
1680	695	825	147.0
2319	0	600	144.0
3000	299	1699	132.0
3004	1099	1499	129.0
...
195905	999	1221	127.0
196141	999	999	130.0
196453	305	1042	130.0
196469	349	899	146.0
196506	149	899	138.0

	total_busy_partners	total_outstanding_orders
170	128.0	261.0
1680	136.0	272.0
2319	138.0	276.0
3000	129.0	256.0
3004	128.0	261.0
...
195905	104.0	251.0
196141	134.0	253.0
196453	141.0	260.0
196469	139.0	272.0
196506	136.0	253.0

[371 rows x 14 columns]

The three operational columns (onshift, busy, outstanding) also contain outliers.

Many rows with null or zero values occur together and show inconsistent behavior.

Data Cleaning

This stage involves detecting and addressing issues like missing values, duplicates, incorrect formatting, or irrelevant data, ensuring the data is accurate, complete, and consistent for analysis or decision-making.

Creating duplicate dataset, as it act as a backup and allows us for experimentation without altering the original data.

```
# creating porter dataframe
pdf = PorterData
```

Fix issues like nulls, wrong formats, and duplicated data before doing any analysis.

```
# dropping rows based on conditions

pdf.drop(pdf[pdf['total_items'] == 411].index, inplace=True) # removing extreme outliers from total items column

pdf.drop(pdf[pdf['subtotal'] == 0].index, inplace=True) # removing missing transaction data

pdf.drop(pdf[pdf['min_item_price'] < 0].index, inplace=True) # removing negative values from column as price is not negative

pdf.drop(pdf[pdf['min_item_price'] > 10000].index, inplace=True) # removing outliers from min item price column

pdf.drop(pdf[pdf['max_item_price'] == 0].index, inplace=True) # removing missing transaction data
```

```

pdf.drop(pdf[pdf['total_onshift_partners'] < 0].index, inplace=True) # removing negative values from column

pdf.drop(pdf[pdf['total_busy_partners'] < 0].index, inplace=True) # removing negative values from column

pdf.drop(pdf[pdf['total_outstanding_orders'] < 0].index, inplace=True) # removing negative values from column

pdf.dropna(subset=['actual_delivery_time'], inplace=True) # Guessing they are canceled orders.

```

Replace nulls in categorical columns (like store type or order protocol) with a generic value like "unknown"

```

# replace all nulls with, 7/unknown category for market_id
pdf['market_id'].fillna(7, inplace=True)

# replace all nulls with 8/unknown category for order_protocol
pdf['order_protocol'].fillna(8, inplace=True)

# fill missing 'store_primary_category' with 'unknown'
pdf['store_primary_category'].fillna('unknown', inplace=True)

```

Convert the created_at column to datetime format so that time-based analysis can be done easily.

```

# convert to datetime
pdf['created_at'] = pd.to_datetime(pdf['created_at'])
pdf['actual_delivery_time'] =
pd.to_datetime(pdf['actual_delivery_time'])

```

Creating features columns like date, hour, minute, dayname, etc, using created_at column
 delivery_duration_min: Calculate time taken for delivery in minutes.

```

# create a new columns for time features
pdf['date'] = pdf['created_at'].dt.date
pdf['hour'] = pdf['created_at'].dt.hour
pdf['minute'] = pdf['created_at'].dt.minute
pdf['dayname'] = pdf['created_at'].dt.day_name() # 0=Monday, 6=Sunday

# calculate delivery duration in minutes
pdf['delivery_duration_min'] = (pdf['actual_delivery_time'] -
pdf['created_at']).dt.total_seconds() / 60

```

available_partners: Calculate using onshift - busy.

partner_load: Orders divided by available partners, helps understand workload.

```

# create available partners = onshift partner - busy partners
pdf['available_partners'] = pdf['total_onshift_partners'] -
pdf['total_busy_partners']
pdf.loc[pdf['available_partners'] < 0, 'available_partners'] = np.nan
# replace negatives with NaN

# create partner load = orders / available partners
pdf['partner_load'] = pdf['total_items'] /
pdf['available_partners'].replace(0, np.nan) # avoid divide by zero

```

Columns like total_onshift_partners, total_busy_partners, and total_outstanding_orders have about 8% missing data.

Deleting these rows would cause data loss, so we choose to fill them using imputation.

Since the data is not normally distributed, and has outliers, median is used to fill missing values for partner-related columns.

```

print(f"Onshift_Mean: {pdf['total_onshift_partners'].mean()}")
print(f"Onshift_Median: {pdf['total_onshift_partners'].median()")

print(f"Busy_Mean: {pdf['total_busy_partners'].mean()}")
print(f"Busy_Median: {pdf['total_busy_partners'].median()}

print(f"Outstanding_Mean: {pdf['total_outstanding_orders'].mean()}")
print(f"Outstanding_Median:
{pdf['total_outstanding_orders'].median()}

Onshift_Mean: 44.82775910488026
Onshift_Median: 37.0
Busy_Mean: 41.7578223470358
Busy_Median: 34.0
Outstanding_Mean: 58.07873206111934
Outstanding_Median: 41.0

# replacing null values with total_onshift_partners.median value
pdf['total_onshift_partners'].fillna(pdf['total_onshift_partners'].median(),
inplace=True)

# replacing null values with total_busy_partners.median value
pdf['total_busy_partners'].fillna(pdf['total_busy_partners'].median(),
inplace=True)

# replacing null values with total_outstanding_orders.median value
pdf['total_outstanding_orders'].fillna(pdf['total_outstanding_orders'].median(),
inplace=True)

# Converting columns datatype
pdf['date'] = pd.to_datetime(pdf['date'])
pdf['market_id'] = pdf['market_id'].astype('object')
pdf['order_protocol'] = pdf['order_protocol'].astype('object')

```

```

# checking information about the data
pdf.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 197139 entries, 0 to 197427
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   market_id        197139 non-null   object  
 1   created_at       197139 non-null   datetime64[ns]
 2   actual_delivery_time 197139 non-null   datetime64[ns]
 3   store_id         197139 non-null   object  
 4   store_primary_category 197139 non-null   object  
 5   order_protocol   197139 non-null   object  
 6   total_items      197139 non-null   int64   
 7   subtotal         197139 non-null   int64   
 8   num_distinct_items 197139 non-null   int64   
 9   min_item_price   197139 non-null   int64   
 10  max_item_price   197139 non-null   int64   
 11  total_onshift_partners 197139 non-null   float64 
 12  total_busy_partners 197139 non-null   float64 
 13  total_outstanding_orders 197139 non-null   float64 
 14  date             197139 non-null   datetime64[ns]
 15  hour             197139 non-null   int64   
 16  minute           197139 non-null   int64   
 17  dayname          197139 non-null   object  
 18  delivery_duration_min 197139 non-null   float64 
 19  available_partners 140578 non-null   float64 
 20  partner_load     104717 non-null   float64 
dtypes: datetime64[ns](3), float64(6), int64(7), object(5)
memory usage: 33.1+ MB

# Check null percentage
pdf.isnull().mean() * 100

market_id           0.000000
created_at          0.000000
actual_delivery_time 0.000000
store_id            0.000000
store_primary_category 0.000000
order_protocol      0.000000
total_items         0.000000
subtotal            0.000000
num_distinct_items 0.000000
min_item_price      0.000000
max_item_price      0.000000
total_onshift_partners 0.000000
total_busy_partners 0.000000
total_outstanding_orders 0.000000
date                0.000000

```

```

hour                      0.000000
minute                     0.000000
dayname                     0.000000
delivery_duration_min      0.000000
available_partners        28.690924
partner_load                46.881642
dtype: float64

# Checking descriptive statistics
pdf.describe()

    total_items      subtotal num_distinct_items
min_item_price \
count 197139.000000 197139.000000      197139.000000
197139.000000
mean      3.193889    2684.840082      2.670887
686.141088
std       2.501450    1822.061562      1.630295
521.023015
min       1.000000     95.000000      1.000000
0.000000
25%       2.000000    1401.000000      1.000000
299.000000
50%       3.000000    2200.000000      2.000000
595.000000
75%       4.000000    3396.000000      3.000000
949.000000
max       84.000000   27100.000000     20.000000
8999.000000

    max_item_price total_onshift_partners total_busy_partners \
count 197139.000000          197139.000000      197139.000000
mean      1159.587915            44.182643      41.118470
std       557.535031            33.141684      30.863559
min       60.000000             0.000000      0.000000
25%       800.000000            19.000000      17.000000
50%       1095.000000            37.000000      34.000000
75%       1395.000000            62.000000      59.000000
max       8999.000000           171.000000     154.000000

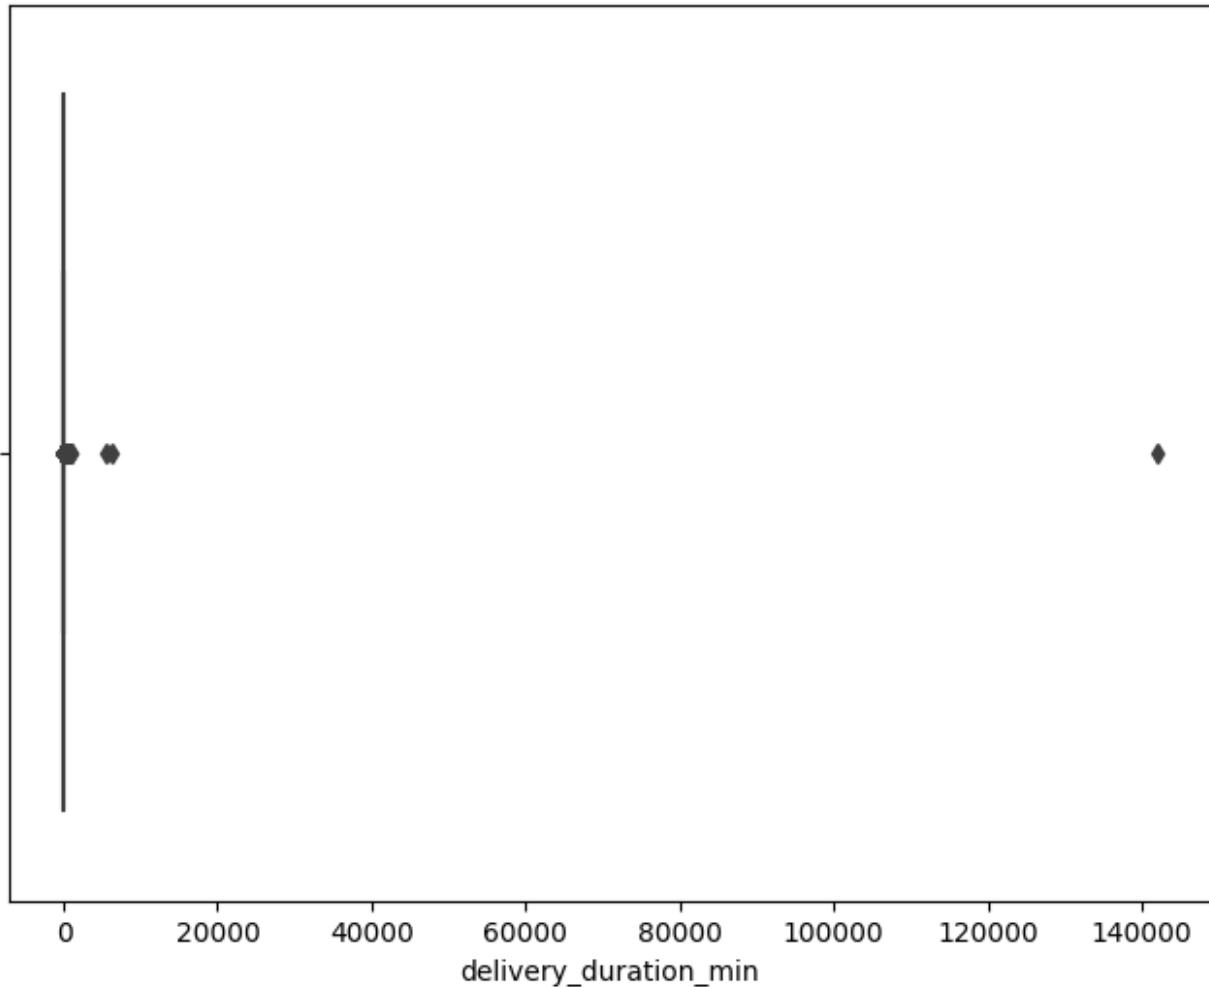
    total_outstanding_orders hour minute \
count 197139.000000 197139.000000 197139.000000
mean      56.671207    8.465950    29.559676
std       50.664727    8.658617    17.277731
min       0.000000     0.000000     0.000000
25%       19.000000    2.000000    15.000000
50%       41.000000    3.000000    29.000000
75%       80.000000   19.000000    45.000000
max      285.000000   23.000000    59.000000

```

	delivery_duration_min	available_partners	partner_load
count	197139.000000	140578.000000	104717.000000
mean	48.470635	6.237555	1.115262
std	320.721991	9.427962	1.608477
min	1.683333	0.000000	0.012500
25%	35.066667	0.000000	0.222222
50%	44.333333	3.000000	0.500000
75%	56.350000	8.000000	1.333333
max	141947.650000	86.000000	48.000000

```
# boxplot for delivery duration column
plt.figure(figsize=(8, 6))
sns.boxplot(x=pdf['delivery_duration_min'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



Column like available_partners and partner_load contain null values, as difference between onshit and busy partners, in some of rows is negative, and replacing this negative to null values, as the result linked to partner_load column.

Delivery Duration contain unrealistic time i.e. outliers, so limiting the rows between 10 minutes and 2 hours

```
# outliers from delivery duration column
((pdf['delivery_duration_min'] < 10) | (pdf['delivery_duration_min'] >
120)).value_counts()

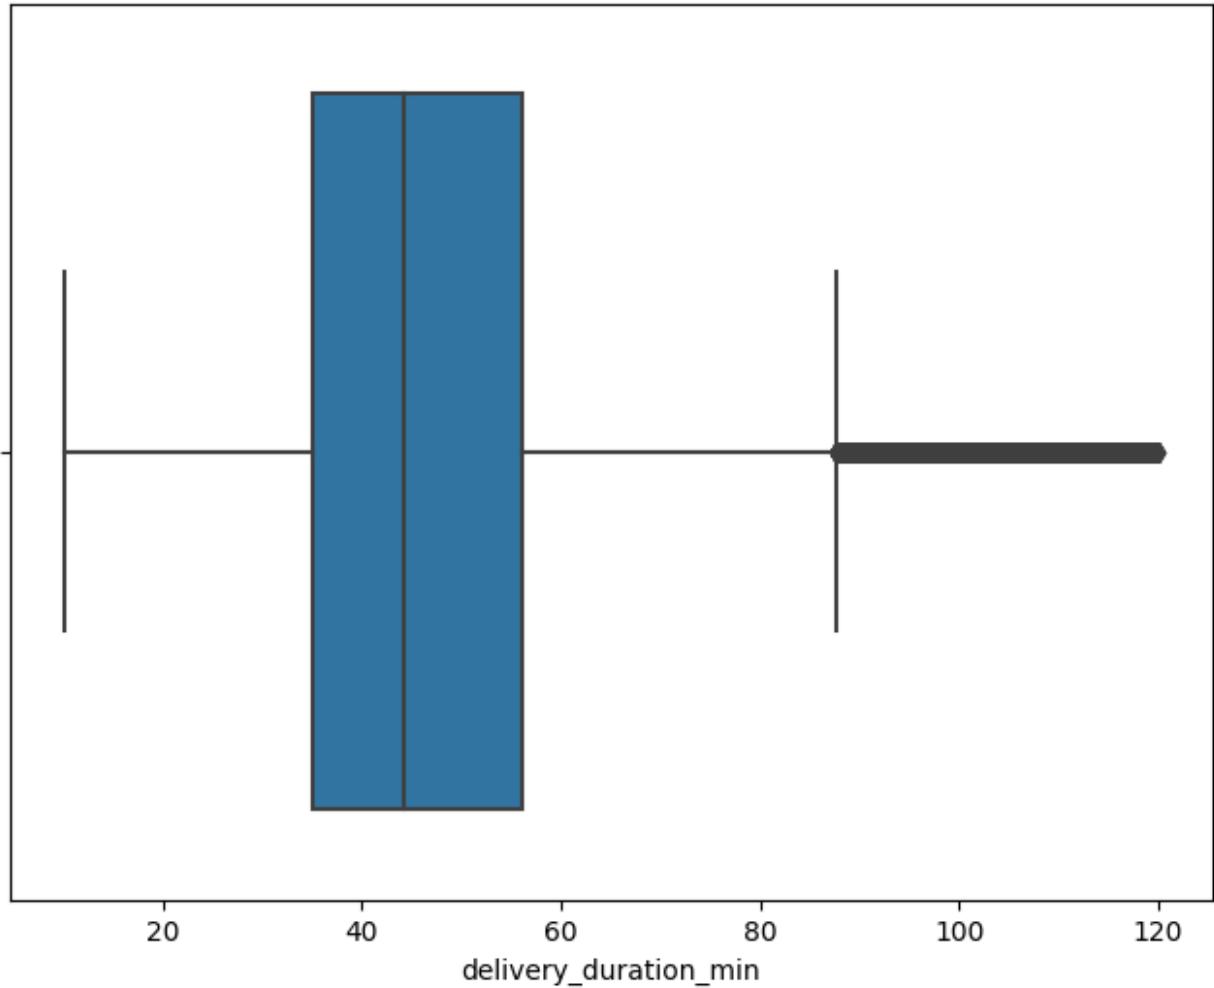
False    196022
True      1117
Name: delivery_duration_min, dtype: int64
```

Only keep records between 10 minutes and 2 hours to remove unrealistic values.

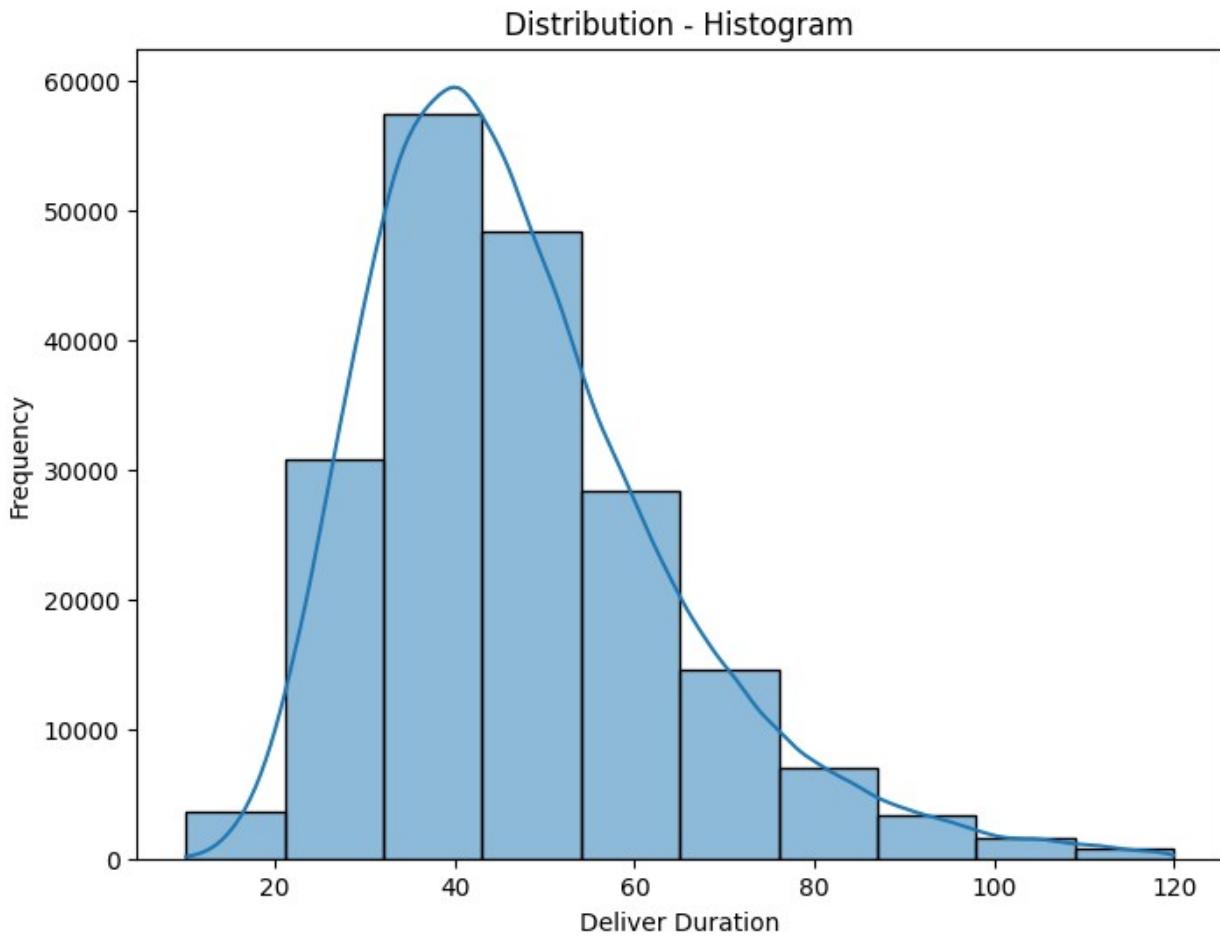
```
# Filter unrealistic delivery times (e.g., <10 mins or >120 mins)
pdf = pdf[(pdf['delivery_duration_min'] >= 10) &
(pdf['delivery_duration_min'] <= 120)]

# boxplot for delivery duration column
plt.figure(figsize=(8, 6))
sns.boxplot(x=pdf['delivery_duration_min'])
plt.title('Distribution - Box Plot')
plt.show()
```

Distribution - Box Plot



```
# histogram for delivery duration column
plt.figure(figsize=(8, 6))
sns.histplot(pdf['delivery_duration_min'], kde=True, bins=10)
plt.title('Distribution - Histogram')
plt.xlabel('Deliver Duration')
plt.ylabel('Frequency')
plt.show()
```



```
pdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 196022 entries, 0 to 197427
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   market_id        196022 non-null   object 
 1   created_at       196022 non-null   datetime64[ns]
 2   actual_delivery_time 196022 non-null   datetime64[ns]
 3   store_id         196022 non-null   object 
 4   store_primary_category 196022 non-null   object 
 5   order_protocol   196022 non-null   object 
 6   total_items      196022 non-null   int64  
 7   subtotal         196022 non-null   int64  
 8   num_distinct_items 196022 non-null   int64  
 9   min_item_price   196022 non-null   int64  
 10  max_item_price   196022 non-null   int64  
 11  total_onshift_partners 196022 non-null   float64
 12  total_busy_partners 196022 non-null   float64
 13  total_outstanding_orders 196022 non-null   float64
```

```

14 date 196022 non-null datetime64[ns]
15 hour 196022 non-null int64
16 minute 196022 non-null int64
17 dayname 196022 non-null object
18 delivery_duration_min 196022 non-null float64
19 available_partners 139870 non-null float64
20 partner_load 104345 non-null float64
dtypes: datetime64[ns](3), float64(6), int64(7), object(5)
memory usage: 32.9+ MB

```

```
pdf.describe()
```

	total_items	subtotal	num_distinct_items
min_item_price \ count	196022.000000	196022.000000	196022.000000
196022.000000			
mean	3.191392	2682.255946	2.669058
686.126685			
std	2.498909	1818.862710	1.627928
520.574802			
min	1.000000	95.000000	1.000000
0.000000			
25%	2.000000	1400.000000	1.000000
299.000000			
50%	3.000000	2200.000000	2.000000
595.000000			
75%	4.000000	3395.000000	3.000000
949.000000			
max	84.000000	26800.000000	20.000000
8999.000000			

	max_item_price	total_onshift_partners	total_busy_partners \
count	196022.000000	196022.000000	196022.000000
mean	1159.187265	44.228796	41.156146
std	556.976319	33.144110	30.866443
min	60.000000	0.000000	0.000000
25%	800.000000	19.000000	17.000000
50%	1095.000000	37.000000	34.000000
75%	1395.000000	62.000000	60.000000
max	8999.000000	171.000000	154.000000

	total_outstanding_orders	hour	minute \
count	196022.000000	196022.000000	196022.000000
mean	56.671312	8.475727	29.558453
std	50.639770	8.661893	17.278957
min	0.000000	0.000000	0.000000
25%	19.000000	2.000000	15.000000
50%	41.000000	3.000000	29.000000
75%	80.000000	19.000000	45.000000
max	285.000000	23.000000	59.000000

	delivery_duration_min	available_partners	partner_load
count	196022.000000	139870.000000	104345.000000
mean	47.100819	6.248552	1.113847
std	16.926042	9.433390	1.606488
min	10.116667	0.000000	0.012500
25%	35.016667	0.000000	0.222222
50%	44.216667	3.000000	0.500000
75%	56.083333	8.000000	1.333333
max	119.933333	86.000000	48.000000

Summarizing Data

```
# earliest and latest timestamps recorded
start_datetime = pdf['created_at'].min()
end_datetime = pdf['created_at'].max()

print(f"Start Date & Time: {start_datetime}")
print(f"End Date & Time: {end_datetime}")

Start Date & Time: 2015-01-21 15:22:03
End Date & Time: 2015-02-18 06:00:44

# total orders
total_orders = pdf.shape[0]
print("Total Orders:", total_orders)

Total Orders: 196022

# total distinct values for market_id, store_id,
# store_primary_category, order_protocol
distinct_counts = {
    'market_id': pdf['market_id'].nunique(),
    'store_id': pdf['store_id'].nunique(),
    'store_primary_category': pdf['store_primary_category'].nunique(),
    'order_protocol': pdf['order_protocol'].nunique()
}

print("Distinct Value Counts:")
for key, value in distinct_counts.items():
    print(f"{key}: {value}")

Distinct Value Counts:
market_id: 7
store_id: 6732
store_primary_category: 75
order_protocol: 8

# average number_of_distinct_items per order
avg_distinct_items = round(pdf['num_distinct_items'].mean())
print(f"Average Distinct Items per Order: {avg_distinct_items:.2f}")
```

```

Average Distinct Items per Order: 3.00

# average orders per week
orders_per_week = pdf['created_at'].dt.to_period('W').value_counts()
avg_orders_per_week = round(orders_per_week.mean())

# average orders per day
orders_per_day = pdf['created_at'].dt.to_period('D').value_counts()
avg_orders_per_day = round(orders_per_day.mean())

# average orders per hour
orders_per_hour = pdf['created_at'].dt.to_period('H').value_counts()
avg_orders_per_hour = round(orders_per_hour.mean())

# average orders per minute
orders_per_minute = pdf['created_at'].dt.to_period('T').value_counts()
avg_orders_per_minute = round(orders_per_minute.mean())

print(f"Average Orders per Week: {avg_orders_per_week:.2f}")
print(f"Average Orders per Day: {avg_orders_per_day:.2f}")
print(f"Average Orders per Hour: {avg_orders_per_hour:.2f}")
print(f"Average Orders per Minute: {avg_orders_per_minute:.4f}")

Average Orders per Week: 39204.00
Average Orders per Day: 6759.00
Average Orders per Hour: 413.00
Average Orders per Minute: 8.0000

# average value for minimum price, maximum price, and subtotal per order
avg_min_price = pdf['min_item_price'].mean()
avg_max_price = pdf['max_item_price'].mean()
avg_subtotal = pdf['subtotal'].mean()

print(f"Average Minimum Price per Order: {avg_min_price:.2f}")
print(f"Average Maximum Price per Order: {avg_max_price:.2f}")
print(f"Average Subtotal Price per Order: {avg_subtotal:.2f}")

Average Minimum Price per Order: 686.13
Average Maximum Price per Order: 1159.19
Average Subtotal Price per Order: 2682.26

# average for onshift / busy partners and outstanding orders
avg_onshift = round(pdf['total_onshift_partners'].mean())
avg_busy = round(pdf['total_busy_partners'].mean())
avg_outstanding = round(pdf['total_outstanding_orders'].mean())

print(f"Average Onshift Partners per Order: {avg_onshift:.2f}")
print(f"Average Busy Partners per Order: {avg_busy:.2f}")
print(f"Average Outstanding Orders per Order: {avg_outstanding:.2f}")

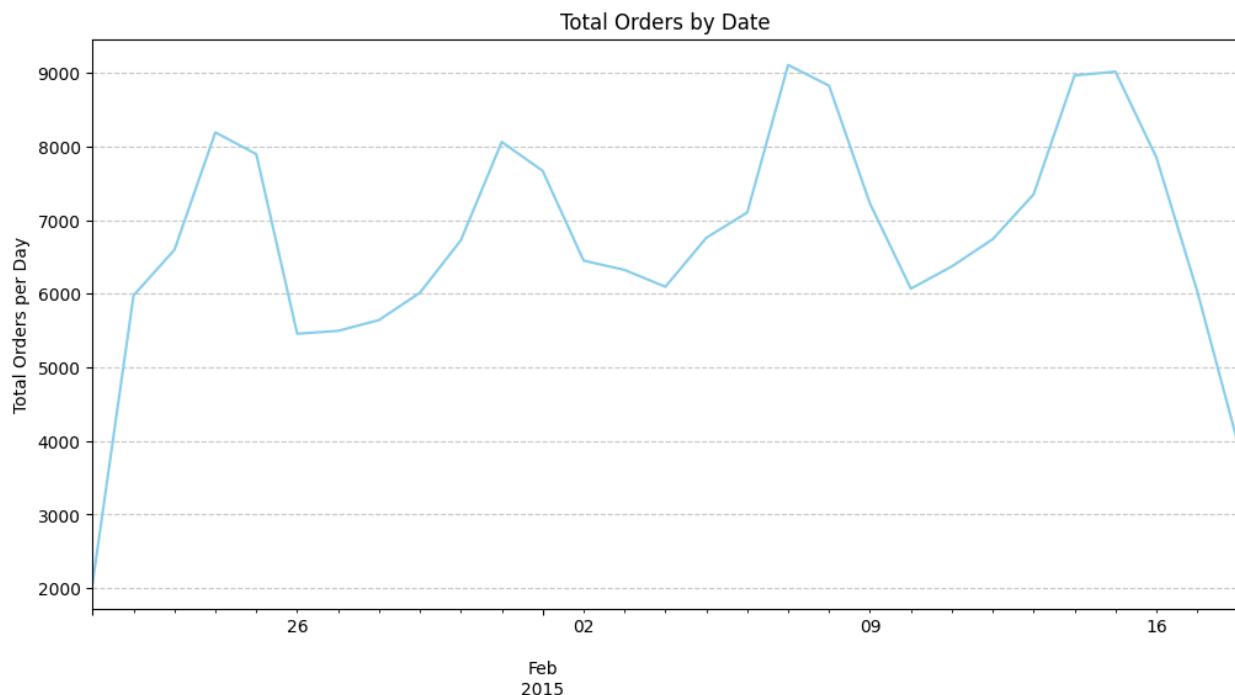
```

```
Average Onshift Partners per Order: 44.00  
Average Busy Partners per Order: 41.00  
Average Outstanding Orders per Order: 57.00
```

Data Visualization

Total Orders by Date

```
# total orders by date  
orders_by_date = pdf.groupby('date').size()  
  
orders_by_date.plot(kind='line', figsize=(12, 6), title='Total Orders  
by Date', color='skyblue')  
plt.xlabel(None)  
plt.ylabel('Total Orders per Day')  
plt.grid(axis='y', linestyle='--', alpha=0.7)  
plt.show()
```



- Daily orders increased steadily over time, showing rising demand for delivery services.
- Orders spike on weekends—especially on Feb 7, Feb 14 (Valentine's Day), and Feb 15, crossing 9,000+ orders.
- After Feb 15, a slight dip is seen—possibly due to reduced weekend activity or external factors.

Average Orders by Day of Week

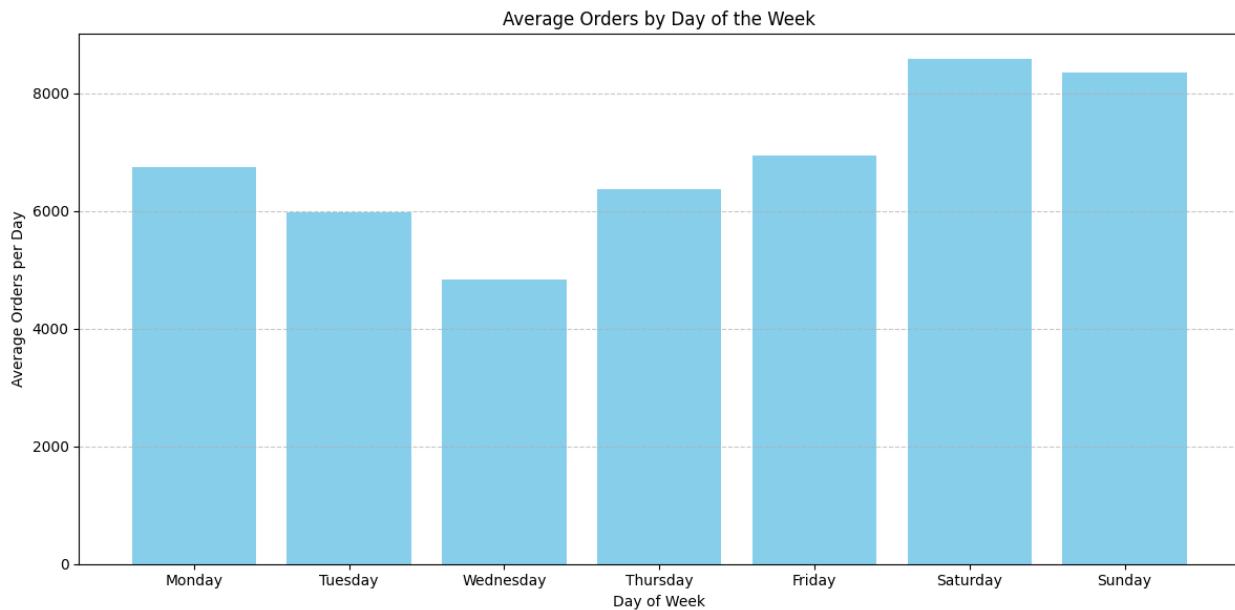
```

# average order by dayname
avg_orders_by_day = (
    pdf.groupby(['dayname', 'date'])
    .size()
    .groupby('dayname')
    .mean()
    .reindex(
        ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
        'Saturday', 'Sunday']
    )
    .reset_index(name='avg_orders')
)

plt.figure(figsize=(12, 6))

plt.bar(avg_orders_by_day['dayname'], avg_orders_by_day['avg_orders'],
color='skyblue')
plt.title('Average Orders by Day of the Week')
plt.xlabel('Day of Week')
plt.ylabel('Average Orders per Day')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



- Saturday (8,580) and Sunday (8,351) see the highest demand, confirming weekends as peak ordering times.
- Friday (6,943) also shows a spike—likely due to weekend planning or celebrations.
- The lowest demand is on Wednesday (4,827).

Average Orders by Hour of Day for Saturday and Sunday

```

# average orders by hour (saturday & sunday)
saturday = pdf[pdf['dayname'] == 'Saturday']
sunday = pdf[pdf['dayname'] == 'Sunday']

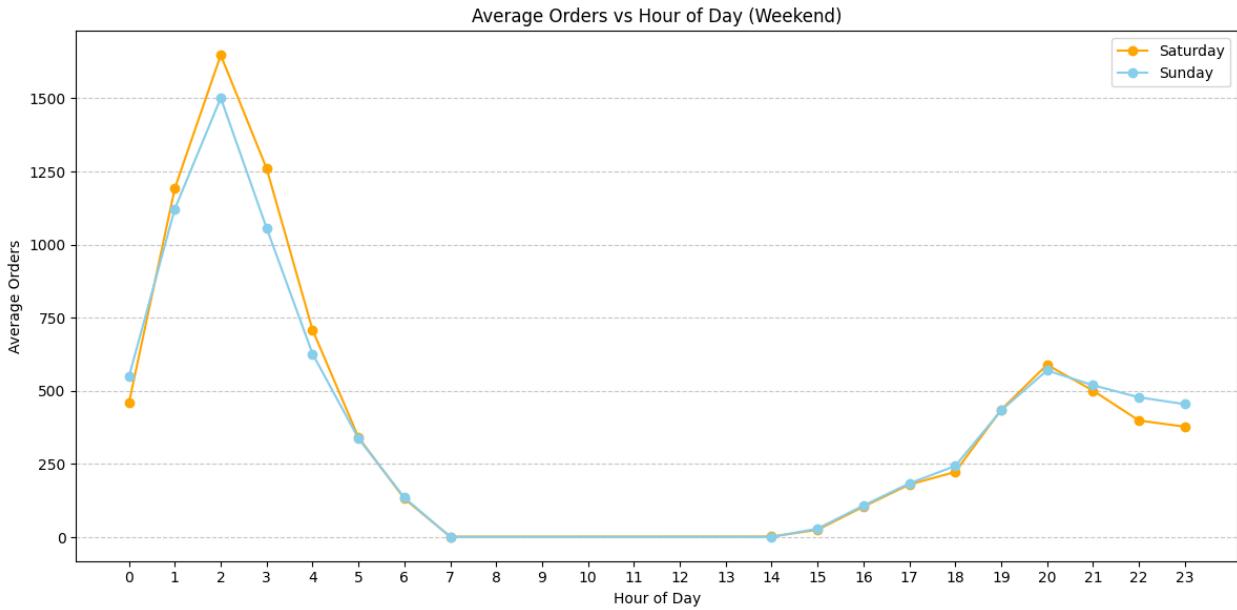
saturday_avg = saturday.groupby(['date',
                                'hour']).size().groupby('hour').mean()
sunday_avg = sunday.groupby(['date',
                            'hour']).size().groupby('hour').mean()

plt.figure(figsize=(12, 6))

plt.plot(saturday_avg.index, saturday_avg.values, label='Saturday',
          color='orange', marker='o')
plt.plot(sunday_avg.index, sunday_avg.values, label='Sunday',
          color='skyblue', marker='o')

plt.title('Average Orders vs Hour of Day (Weekend) ')
plt.xlabel('Hour of Day')
plt.ylabel('Average Orders')
plt.xticks(range(0, 24))
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend()
plt.tight_layout()
plt.show()

```



- Peak hours for both days fall between 1 AM – 3 AM, again reinforcing the late-night hunger trend.
- A minor evening spike from 5 PM–9 PM, especially on Sunday evenings, shows last-minute weekend treats.

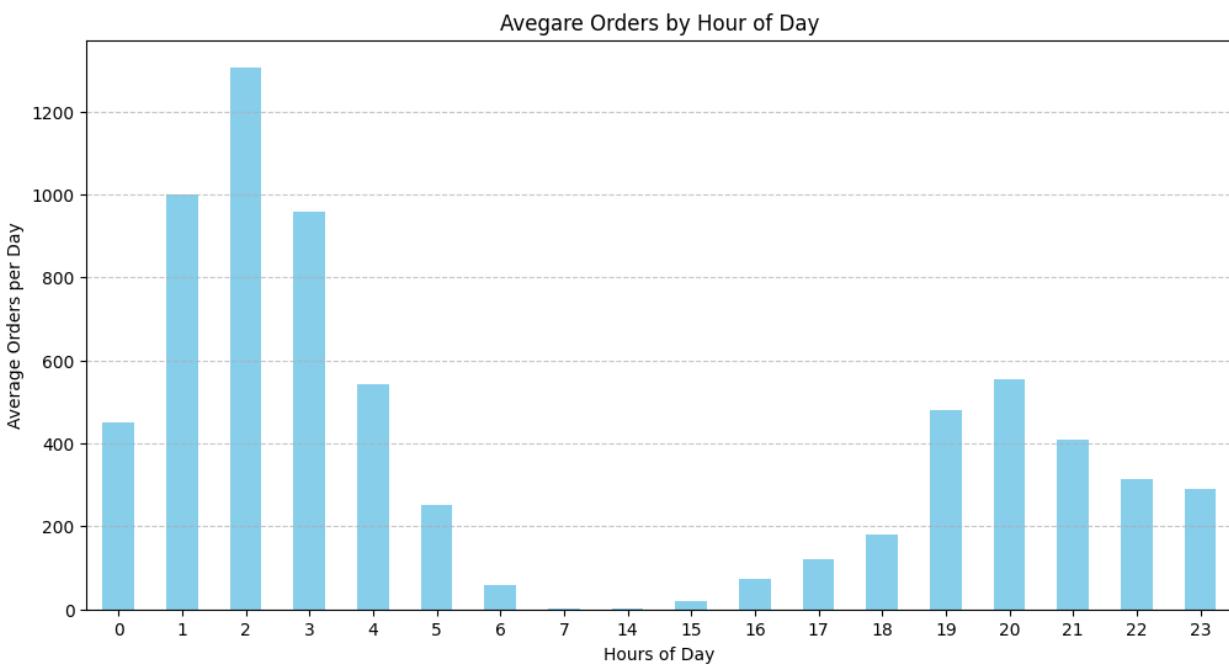
Average Orders by Hour of Day

```

# average orders by hour of day
orders_by_date_hour = pdf.groupby(['date',
'hour']).size().reset_index(name='orders')
avg_orders_by_hour = orders_by_date_hour.groupby('hour')
['orders'].mean()

avg_orders_by_hour.plot(kind='bar', figsize=(12,6), title='Avegare
Orders by Hour of Day', color='skyblue')
plt.xlabel('Hours of Day')
plt.ylabel('Average Orders per Day')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=0)
plt.show()

```



- Highest demand is during early hours (1 AM–3 AM) — peaking around 2 AM (1,306).
- Possibly due to late-night cravings or post-party orders.
- Another small pickup starts around 5 PM–8 PM, suggesting early dinner or evening snacks.
- Very few orders from 6 AM to 2 PM, with a dip to almost zero between 7 AM–2 PM.

Average Orders by Store Primary Category

```

# top 10 store category by average orders
daily_orders_by_category = pdf.groupby(['store_primary_category',
'date']).size().reset_index(name='orders')
avg_orders_by_category =
daily_orders_by_category.groupby('store_primary_category')
['orders'].mean().reset_index()

```

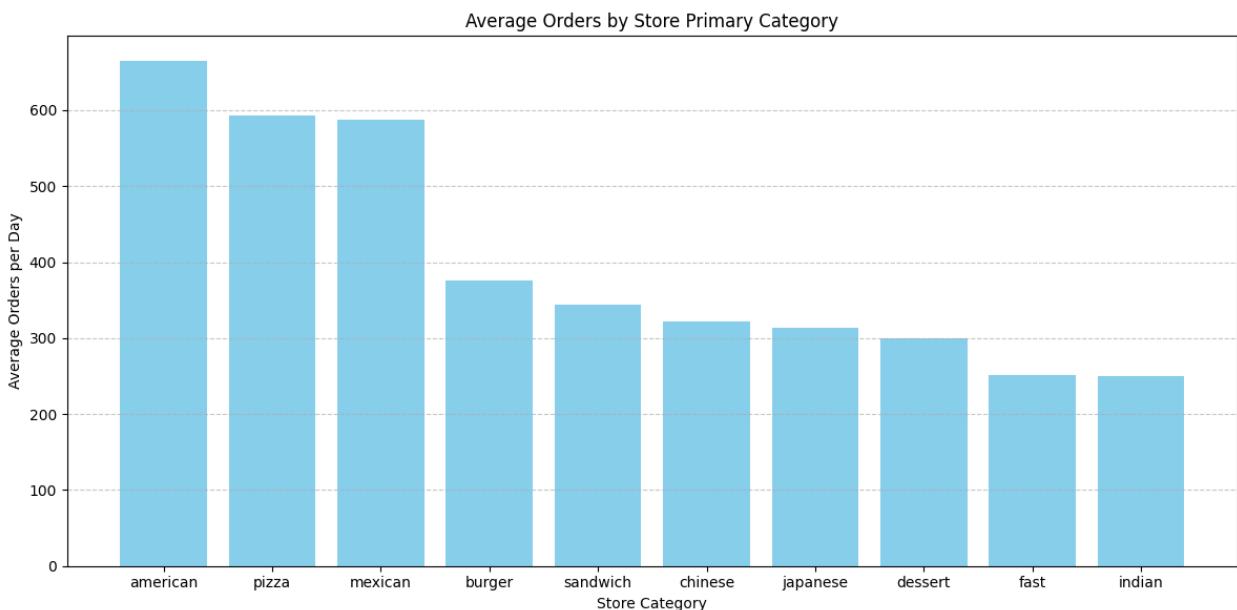
```

avg_orders_by_category =
avg_orders_by_category.sort_values(by='orders',
ascending=False).head(10)

plt.figure(figsize=(12, 6))

plt.bar(avg_orders_by_category['store_primary_category'],
avg_orders_by_category['orders'], color='skyblue')
plt.title('Average Orders by Store Primary Category')
plt.xlabel("Store Category")
plt.ylabel("Average Orders per Day")
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```



- American cuisine leads the way with the highest average number of orders (664), suggesting a strong customer preference for this food category.
- Pizza and Mexican foods follow closely, with 592 and 586 average orders respectively, indicating high demand for comfort and fast foods.
- Burgers, sandwiches, and Chinese dishes round out the mid-tier, showing consistent popularity across different customer segments.
- Indian food is in the 10th position, with a relatively lower average order volume (250), suggesting potential growth opportunity in this category depending on location.

Average Orders by Market Id

```

# average order by market id
avg_orders_by_market =

```

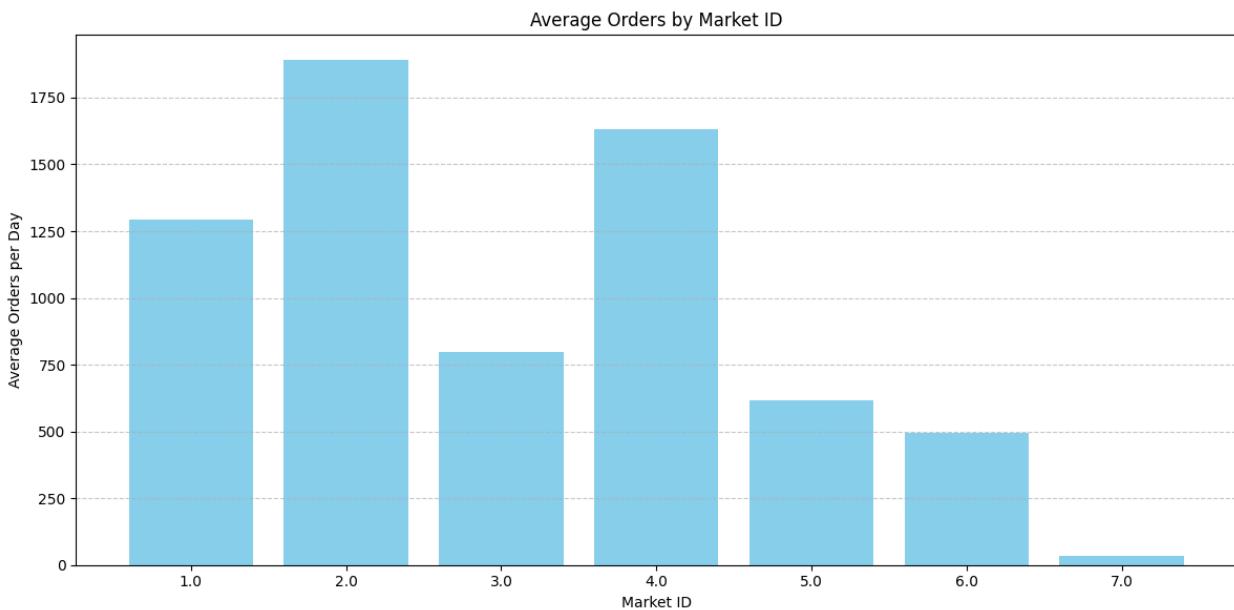
```

pdf.groupby(['market_id', 'date'])
.size()
.groupby('market_id')
.mean()
.reset_index(name='avg_orders')
)

plt.figure(figsize=(12, 6))

plt.bar(avg_orders_by_market['market_id'].astype(str),
avg_orders_by_market['avg_orders'], color='skyblue')
plt.title('Average Orders by Market ID')
plt.xlabel('Market ID')
plt.ylabel('Average Orders per Day')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



- Market 2 (1,890) and Market 4 (1,632) are the busiest.
- Market 6, comparatively, has low demand (495)—indicating a new, niche, or underserved location.

Average Orders by Order Protocol

```

# average orders by order protocol
avg_orders_by_protocol = (
    pdf.groupby(['order_protocol', 'date'])
.size()
.groupby('order_protocol')

```

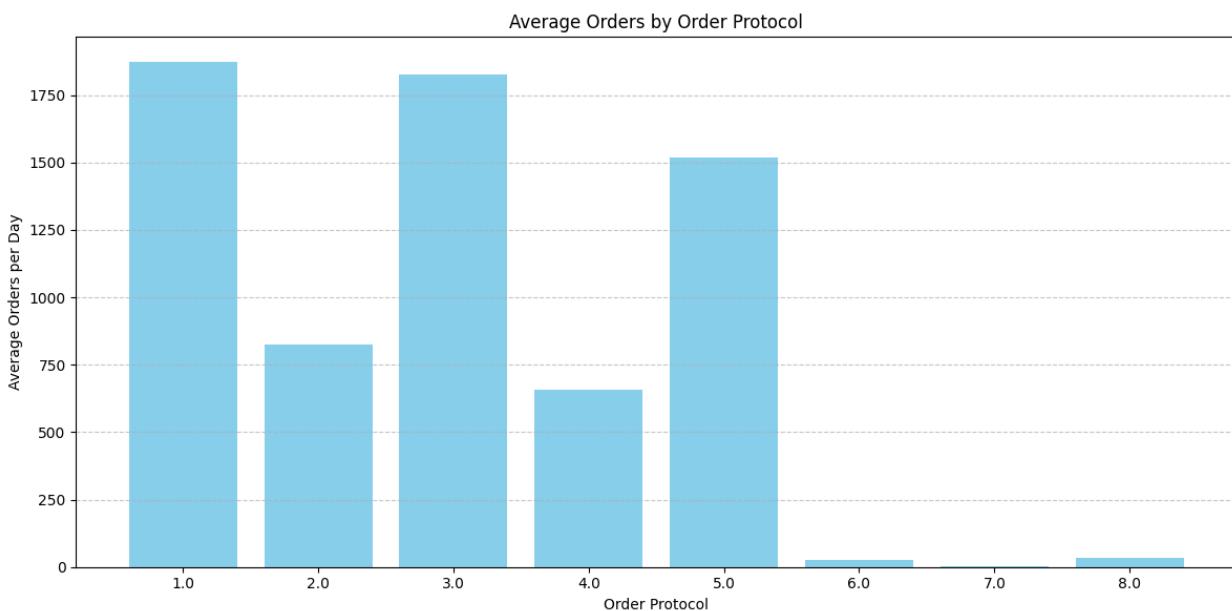
```

    .mean()
    .reset_index(name='avg_orders')
)

plt.figure(figsize=(12, 6))

plt.bar(avg_orders_by_protocol['order_protocol'].astype(str),
        avg_orders_by_protocol['avg_orders'], color='skyblue')
plt.title('Average Orders by Order Protocol')
plt.xlabel('Order Protocol')
plt.ylabel('Average Orders per Day')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



- Protocols 1.0 (1,872) and 3.0 (1,824) are most used — likely mobile app and web orders.
- Lower protocols like 6 and 7 show minimal use — could be legacy systems or internal channels.

Average Delivery Duration by Date

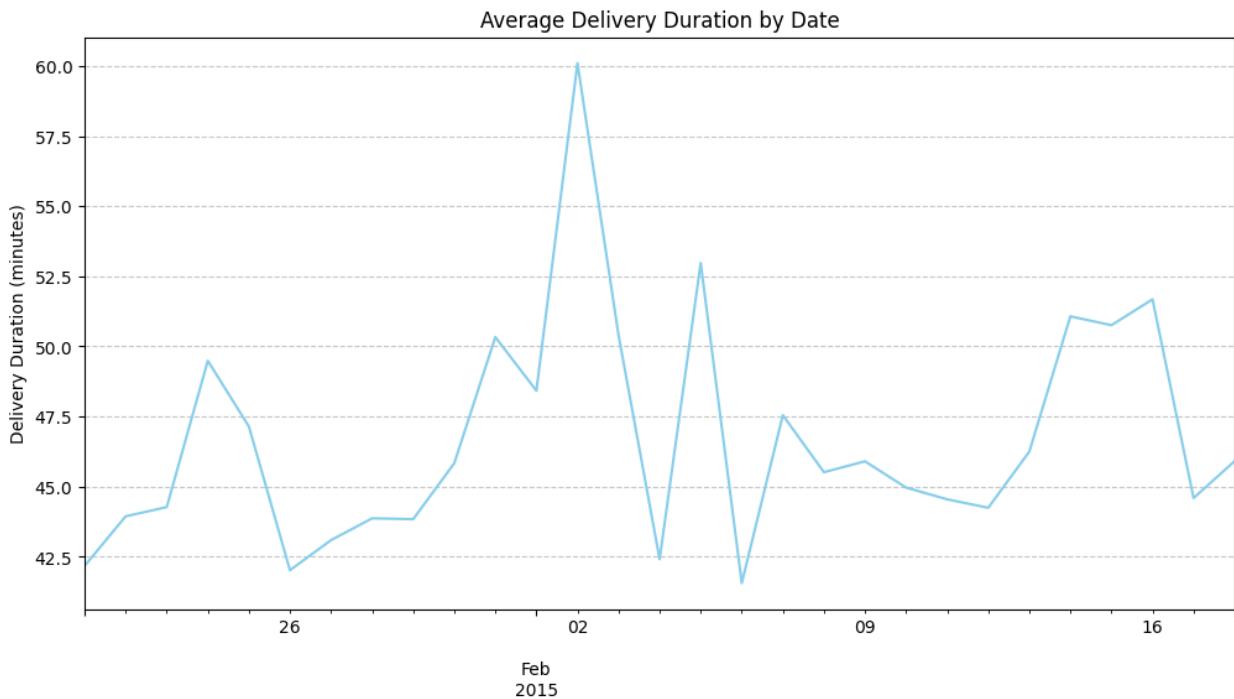
```

# average delivery duration by date
avg_duration_by_date = pdf.groupby('date')
['delivery_duration_min'].mean()

avg_duration_by_date.plot(kind='line', figsize=(12, 6), title='Average
Delivery Duration by Date', color='skyblue')
plt.xlabel(None)
plt.ylabel('Delivery Duration (minutes)')

```

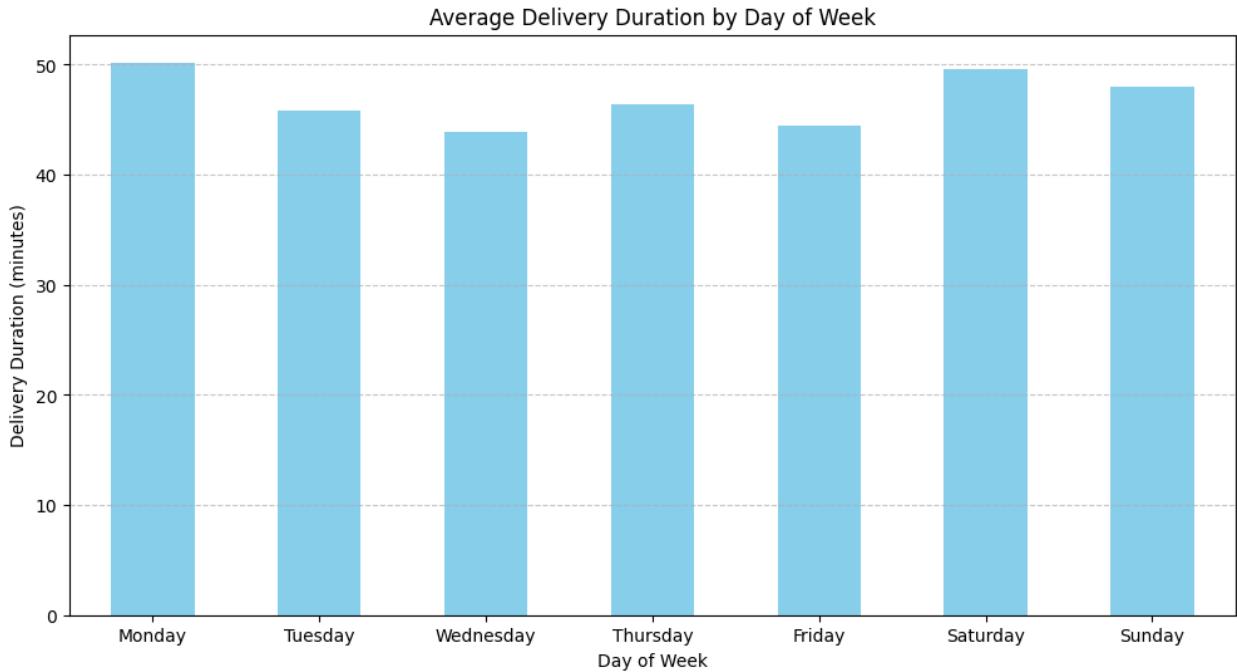
```
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



- There's a noticeable spike in delivery duration around Feb 2, 2015 (60 mins).
 - Weekends like Jan 31 and Feb 14-16 also show higher delivery times, possibly due to higher demand..

Average Delivery Duration by Day of Week

```
# average duration by dayname
avg_duration_by_day = (
    pdf.groupby('dayname')['delivery_duration_min']
    .mean()
    .reindex(
        ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
        'Saturday', 'Sunday']
    )
)
avg_duration_by_day.plot(kind='bar', figsize=(12, 6), title='Average
Delivery Duration by Day of Week', color='skyblue')
plt.xlabel('Day of Week')
plt.ylabel('Delivery Duration (minutes)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

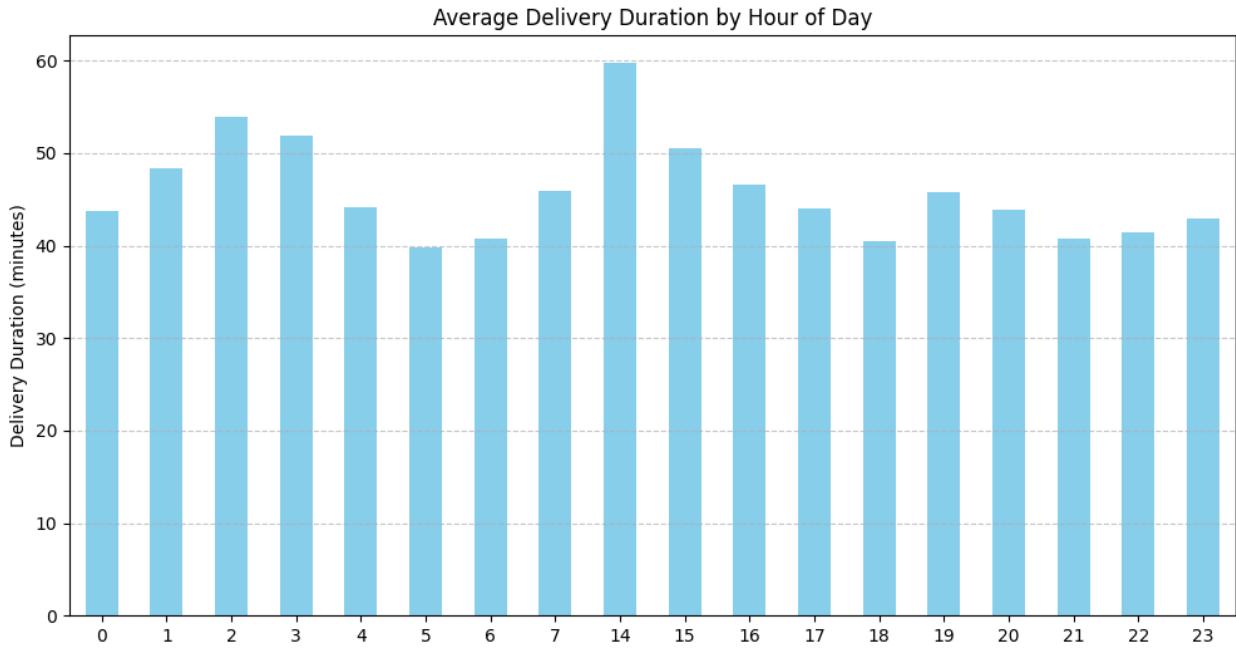


- Monday (50.2 mins) and Saturday (49.6 mins) show the highest average durations.
- Wednesday (43.9 mins) is the most efficient day on average.
- Consider allocating more resources on Mondays and Saturdays.

Average Delivery Duration by Hour of Day

```
# average duration by hour
avg_duration_by_hour = (
    pdf.groupby('hour')['delivery_duration_min']
        .mean()
)

avg_duration_by_hour.plot(kind='bar', figsize=(12,6), title='Average
Delivery Duration by Hour of Day', color='skyblue')
plt.xlabel(None)
plt.ylabel('Delivery Duration (minutes)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

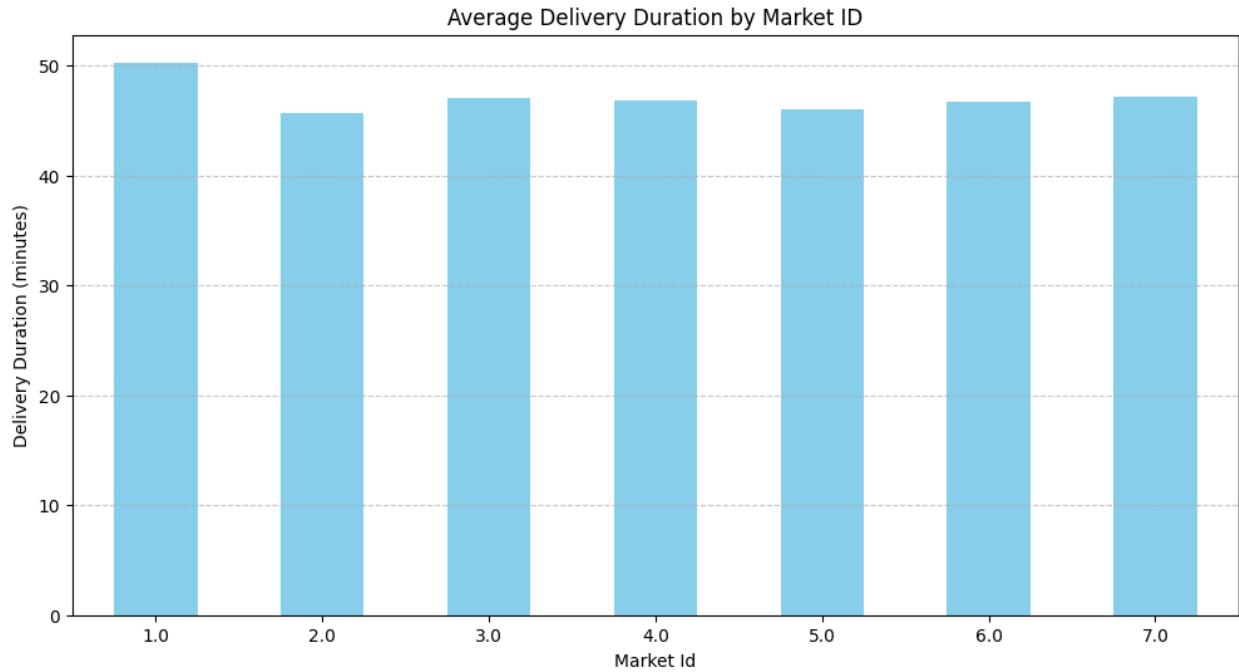


- 2 AM (54 mins) and 3 AM (52 mins) have the longest delivery durations.
- 5 AM – 6 AM (40 mins) and 9 PM – 11 PM (40-43 mins) are the most efficient periods.
- Resource allocation or order demand may vary significantly by hour—optimize late-night staffing.

Average Delivery Duration by Market Id

```
# average duration by market
avg_duration_by_market = pdf.groupby('market_id')
['delivery_duration_min'].mean()

avg_duration_by_market.plot(kind='bar', figsize=(12,6), title='Average
Delivery Duration by Market ID', color='skyblue')
plt.xlabel('Market Id')
plt.ylabel('Delivery Duration (minutes)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

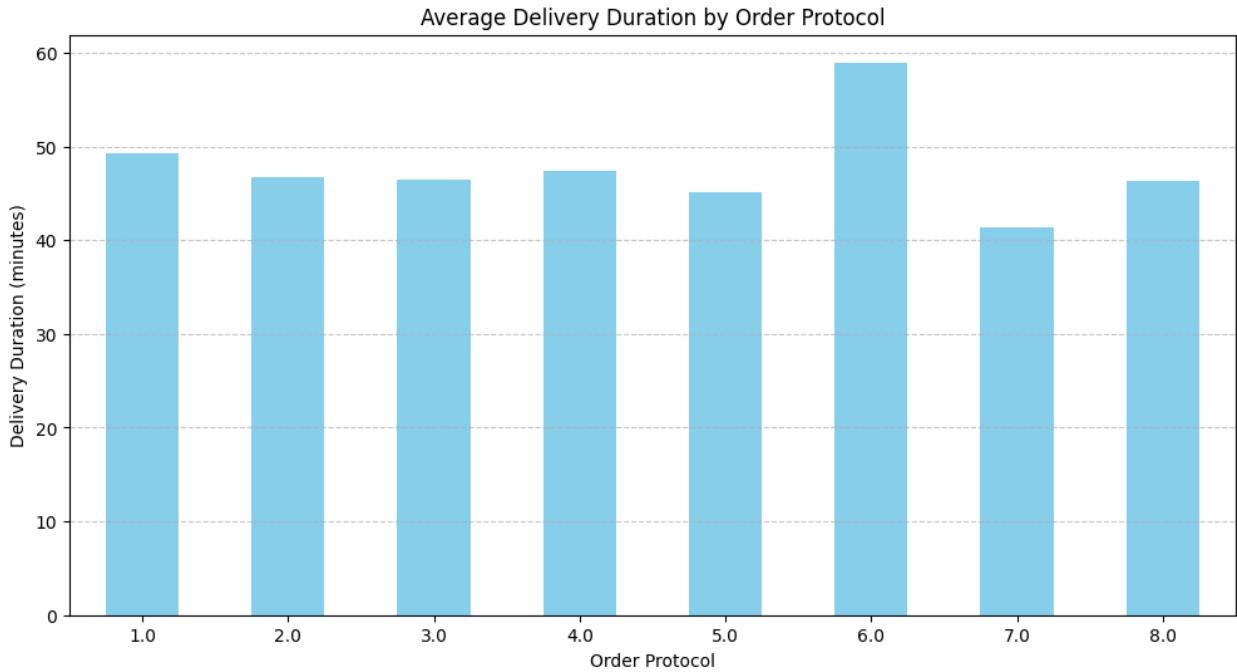


- Market 1 has the longest average time (50.3 mins).
- Market 2 performs the best (45.7 mins).
- Investigate local market issues or traffic conditions in Market 1.

Average Delivery Duration by Order Protocol

```
# average duration by protocol
avg_duration_by_protocol = pdf.groupby('order_protocol')[['delivery_duration_min']].mean()

avg_duration_by_protocol.plot(kind='bar', figsize=(12,6),
title='Average Delivery Duration by Order Protocol', color='skyblue')
plt.xlabel('Order Protocol')
plt.ylabel('Delivery Duration (minutes)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

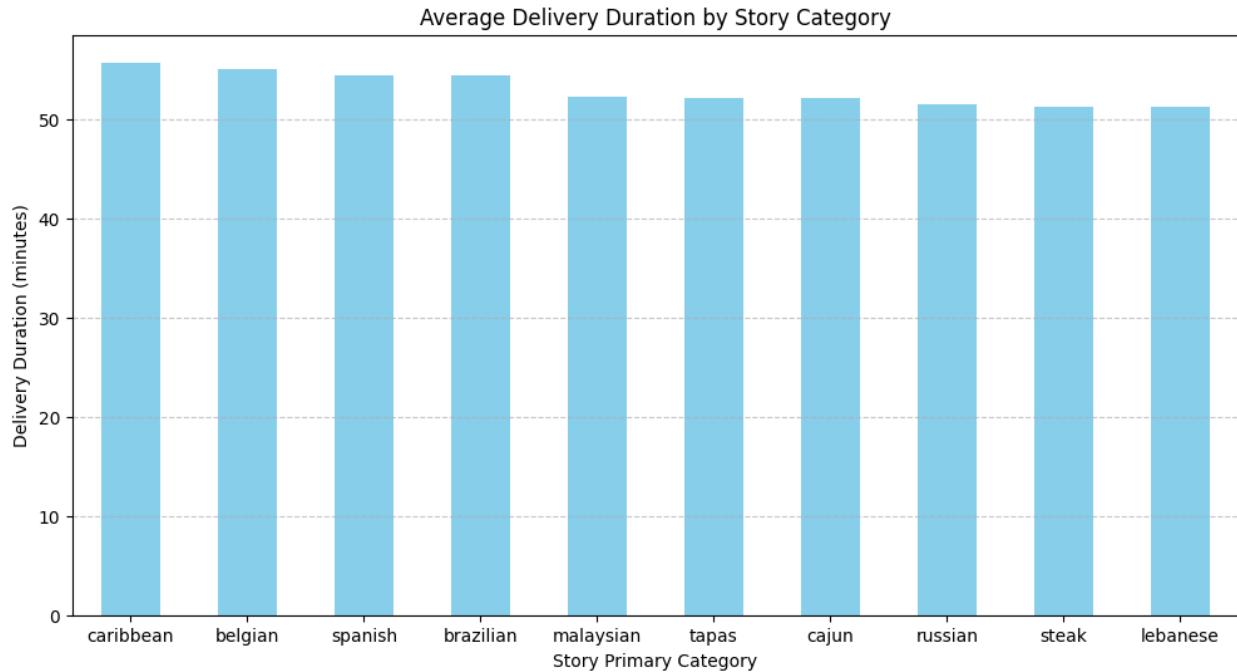


- Protocol 6 has an exceptionally high delivery time (59 mins).
- Protocol 7 is the best performer (41 mins).
- Analyze process inefficiencies or routing strategies for Protocol 6.

Average Delivery Duration by Store Primary Category

```
# average duation by store category
avg_duration_by_category = pdf.groupby('store_primary_category')[['delivery_duration_min']].mean().sort_values(ascending=False).head(10)

avg_duration_by_category.plot(kind='bar', figsize=(12,6),
title='Average Delivery Duration by Story Category', color='skyblue')
plt.xlabel('Story Primary Category')
plt.ylabel('Delivery Duration (minutes)')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

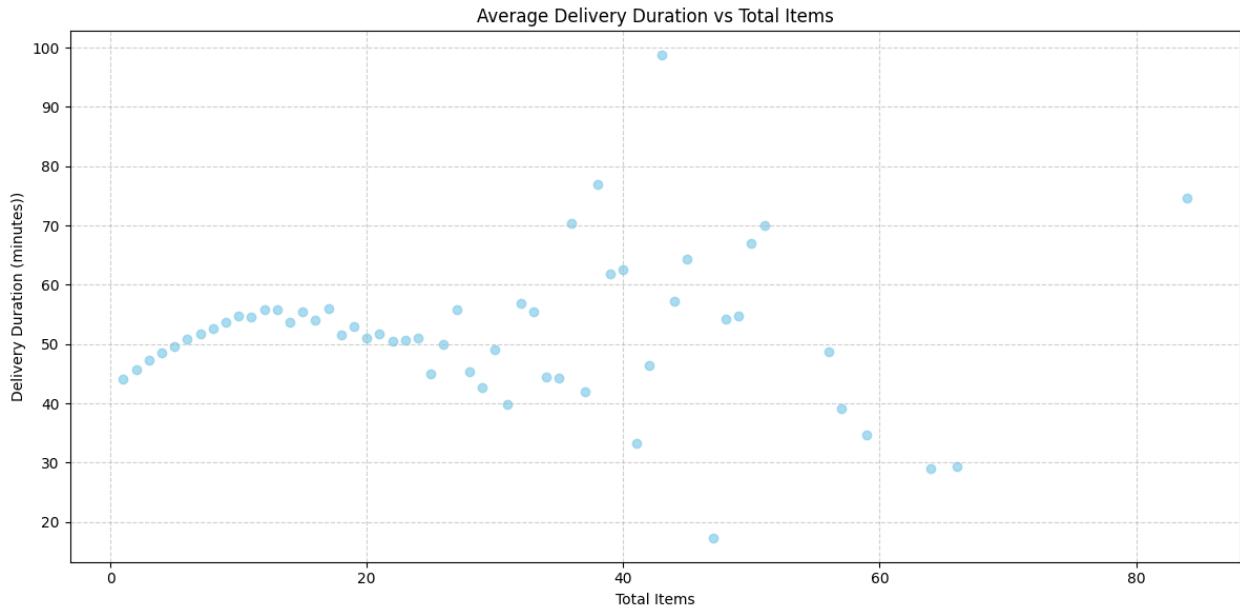


- Categories like Caribbean, Belgian, Spanish, and Brazilian consistently show higher average delivery times (~54–56 mins).
- These cuisines may require more prep time. Consider prepping strategies or informing users about expected delays.

Average Delivery Duration by Total Items

```
# average duration by items
avg_duration_by_items = pdf.groupby('total_items')
['delivery_duration_min'].mean().reset_index()

plt.figure(figsize=(12, 6))
plt.scatter(avg_duration_by_items['total_items'],
avg_duration_by_items['delivery_duration_min'], color='skyblue',
alpha=0.7)
plt.title('Average Delivery Duration vs Total Items')
plt.xlabel('Total Items')
plt.ylabel('Delivery Duration (minutes)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

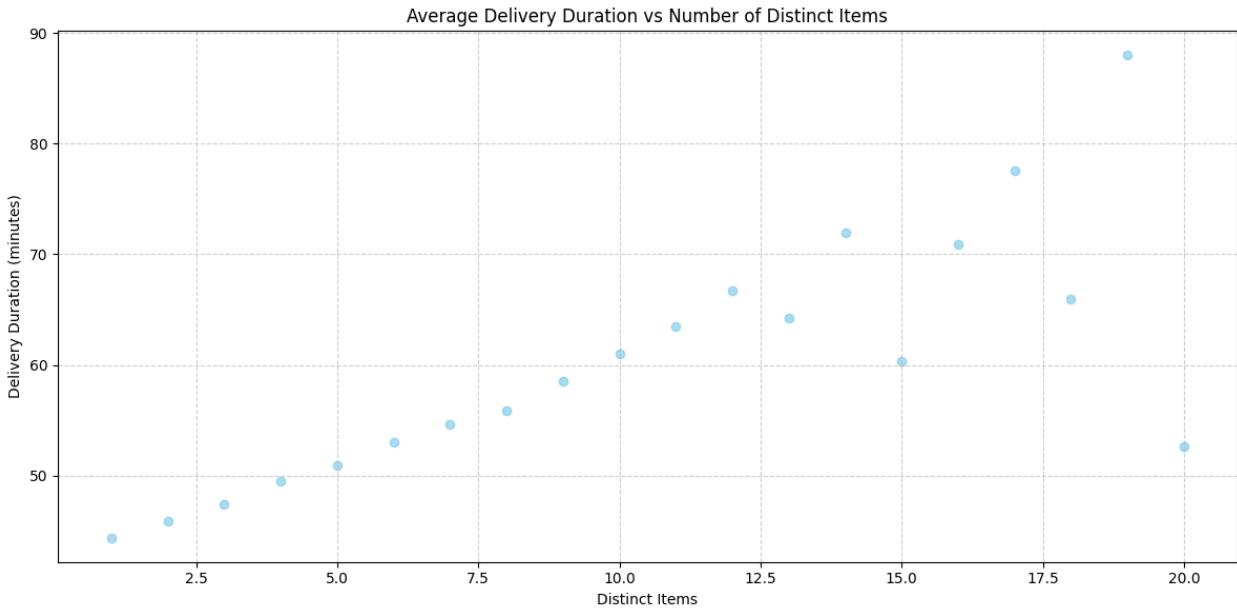


- There is a clear increasing trend in delivery time as the number of items increases.
- Orders with 36+ items see times reaching up to 98 mins.
- However, some outliers (e.g., 47 items = 17 mins) suggest inconsistencies.
- Bulk orders impact delivery time — consider batch processing or special handling for large orders.

Average Delivery Duration by Number of Distinct Items

```
# average duration by distinct items
avg_duration_by_dict_items = pdf.groupby('num_distinct_items')
['delivery_duration_min'].mean().reset_index()

plt.figure(figsize=(12, 6))
plt.scatter(avg_duration_by_dict_items['num_distinct_items'],
avg_duration_by_dict_items['delivery_duration_min'], color='skyblue',
alpha=0.7)
plt.title('Average Delivery Duration vs Number of Distinct Items')
plt.xlabel('Distinct Items')
plt.ylabel('Delivery Duration (minutes)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

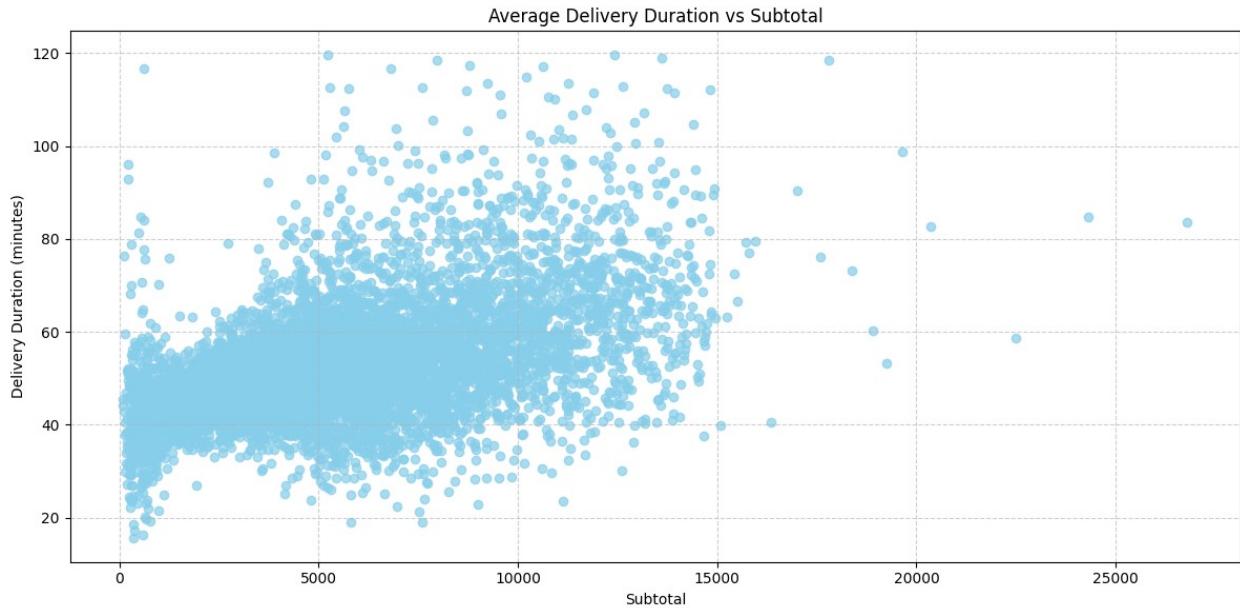


- There's a clear upward trend — as the number of distinct items in an order increases, average delivery duration also increases.
- This is expected, as more variety generally implies more preparation and packaging time

Average Delivery Duration by Subtotal

```
# average duration by subtotal
avg_duration_by_subtotal = pdf.groupby('subtotal')[['delivery_duration_min']].mean().reset_index()

plt.figure(figsize=(12, 6))
plt.scatter(avg_duration_by_subtotal['subtotal'],
            avg_duration_by_subtotal['delivery_duration_min'], color='skyblue',
            alpha=0.7)
plt.title('Average Delivery Duration vs Subtotal')
plt.xlabel('Subtotal')
plt.ylabel('Delivery Duration (minutes)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

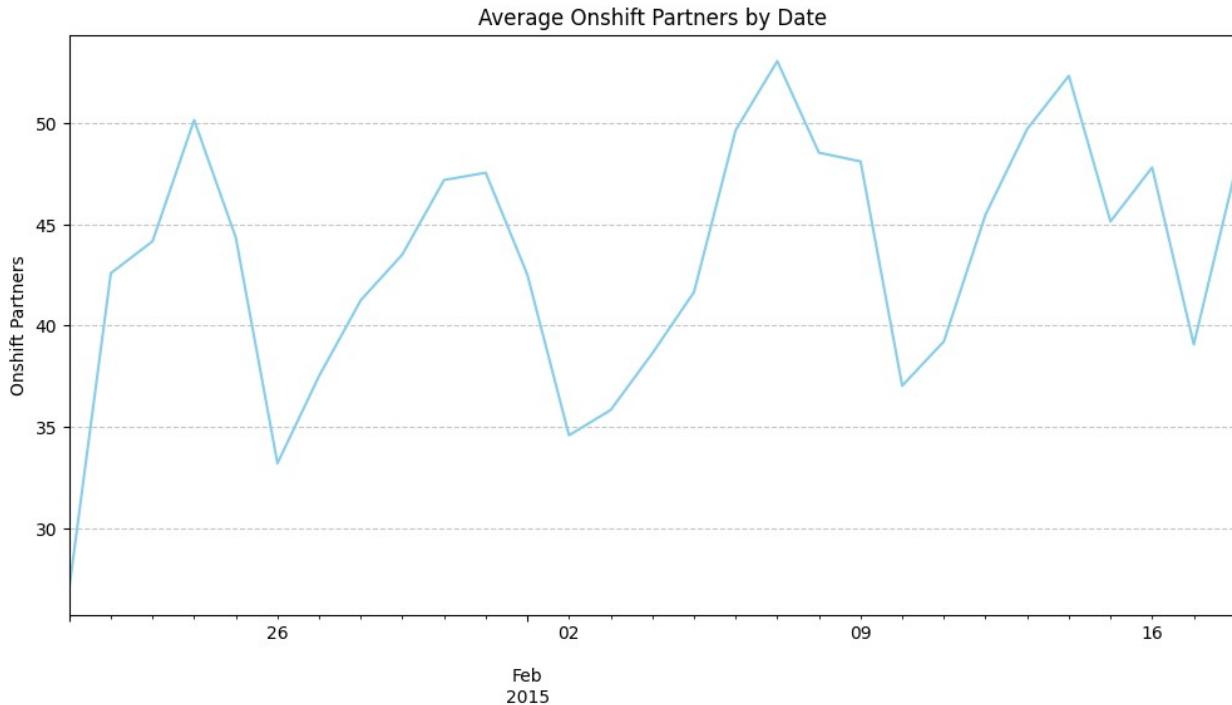


- There is a general upward trend, indicating that as the subtotal increases, the average delivery duration also tends to increase. This could be due to larger or more complex orders requiring more preparation and delivery time.
- Some points deviate significantly from the trend. For example, orders with high subtotals but relatively low delivery durations suggest inconsistencies or exceptional cases (e.g., expedited delivery).
- The data appears to form clusters at certain subtotal ranges, possibly reflecting common order sizes or pricing tiers.
- Orders with subtotals above 20,000 show longer delivery durations, often exceeding 80 minutes. These could represent bulk or special orders.
- Orders with subtotals below 1,000 generally have shorter delivery durations, often under 50 minutes, likely due to their simplicity.

Average Onshift Partners by Date

```
# average onshift partners by date
avg_onshift_by_date = pdf.groupby('date')
['total_onshift_partners'].mean()

avg_onshift_by_date.plot(kind='line', figsize=(12,6), title='Average
Onshift Partners by Date', color='skyblue')
plt.xlabel(None)
plt.ylabel('Onshift Partners')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



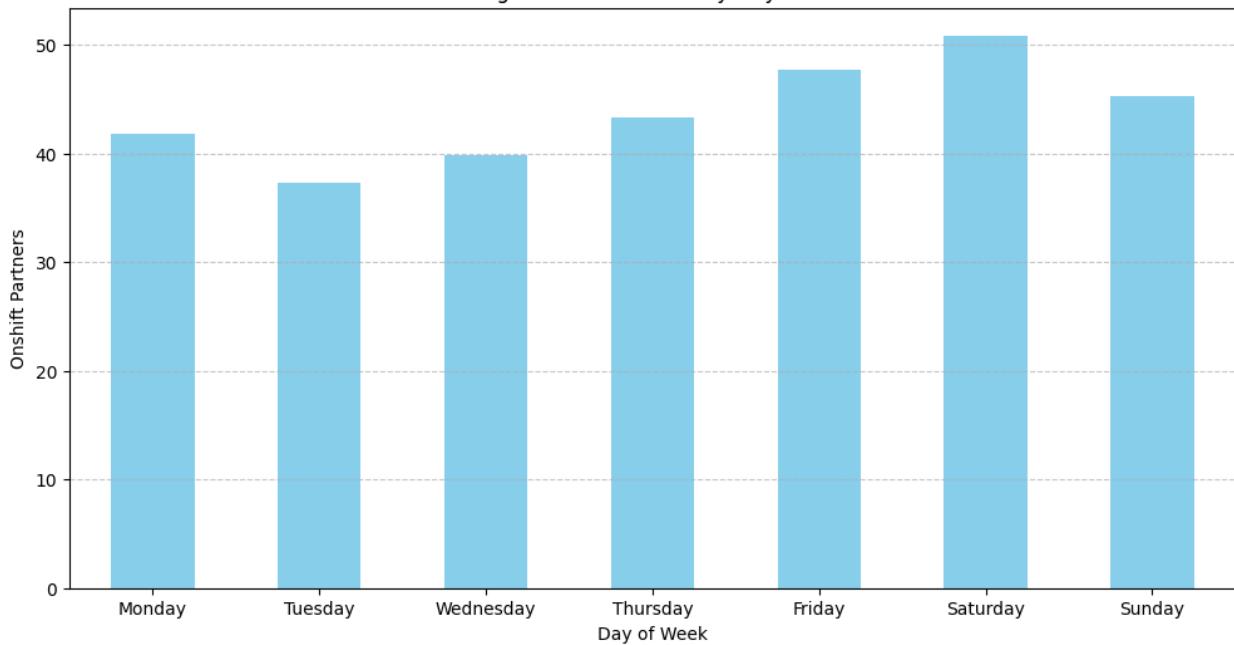
- There is a steady rise in the number of onshift partners from Jan 21 to Feb 18, starting from around 27 and peaking over 53.
- The highest availability occurred during weekends and early February, suggesting stronger supply alignment with anticipated demand.
- Noticeable dips on Mondays and early weekdays may indicate lower partner availability or operational shifts.

Average Onshift Partners by Day of Week

```
# average onshift partners by day
avg_onshift_by_day = (
    pdf.groupby('dayname')['total_onshift_partners']
    .mean()
    .reindex(
        ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
        'Saturday', 'Sunday']
    )
)

avg_onshift_by_day.plot(kind='bar', figsize=(12, 6), title='Average
Onshift Partners by Day of Week', color='skyblue')
plt.xlabel('Day of Week')
plt.ylabel('Onshift Partners')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

Average Onshift Partners by Day of Week



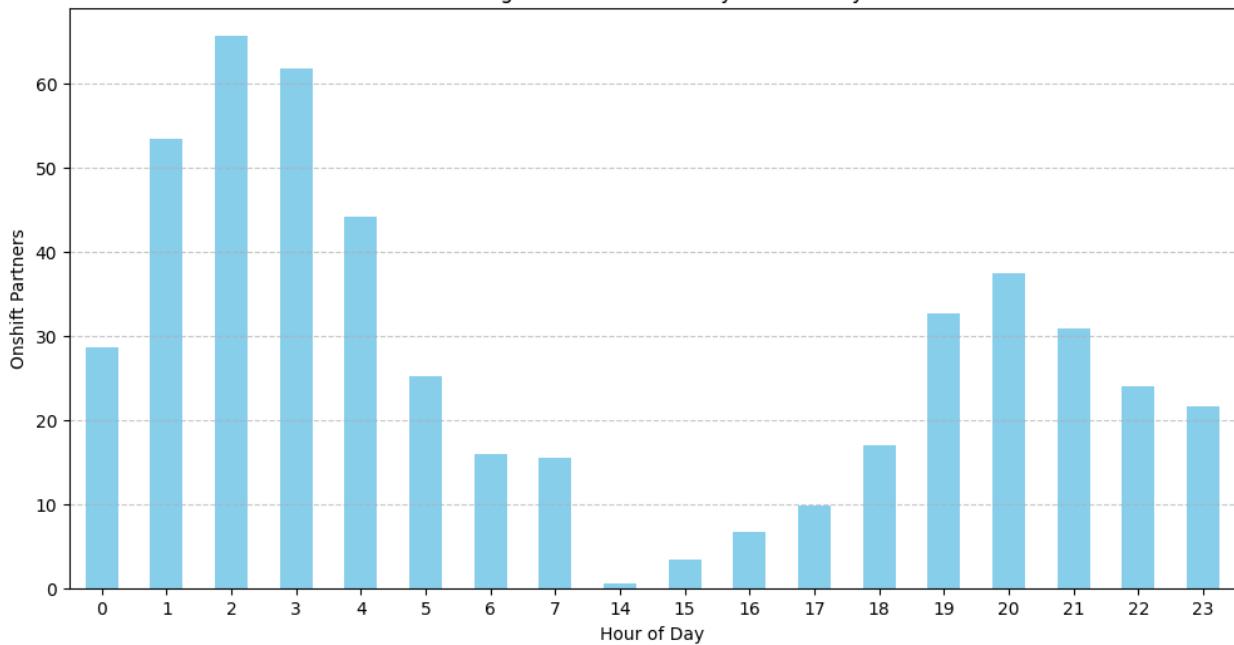
- Saturday (56) and Friday (48) had the highest onshift partner averages, aligning with peak demand days.
- The lowest average is on Tuesday (37), suggesting potential under-capacity midweek.
- There's a clear increase in capacity toward the weekend, likely to handle higher delivery volumes.

Average Onshift Partners by Hour of Day

```
# average onshift partners by hour
avg_onshift_by_hour = (pdf.groupby('hour')
['total_onshift_partners'].mean())

avg_onshift_by_hour.plot(kind='bar', figsize=(12,6), title='Average
Onshift Partners by Hour of Day', color='skyblue')
plt.xlabel('Hour of Day')
plt.ylabel('Onshift Partners')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=0)
plt.show()
```

Average Onshift Partners by Hour of Day

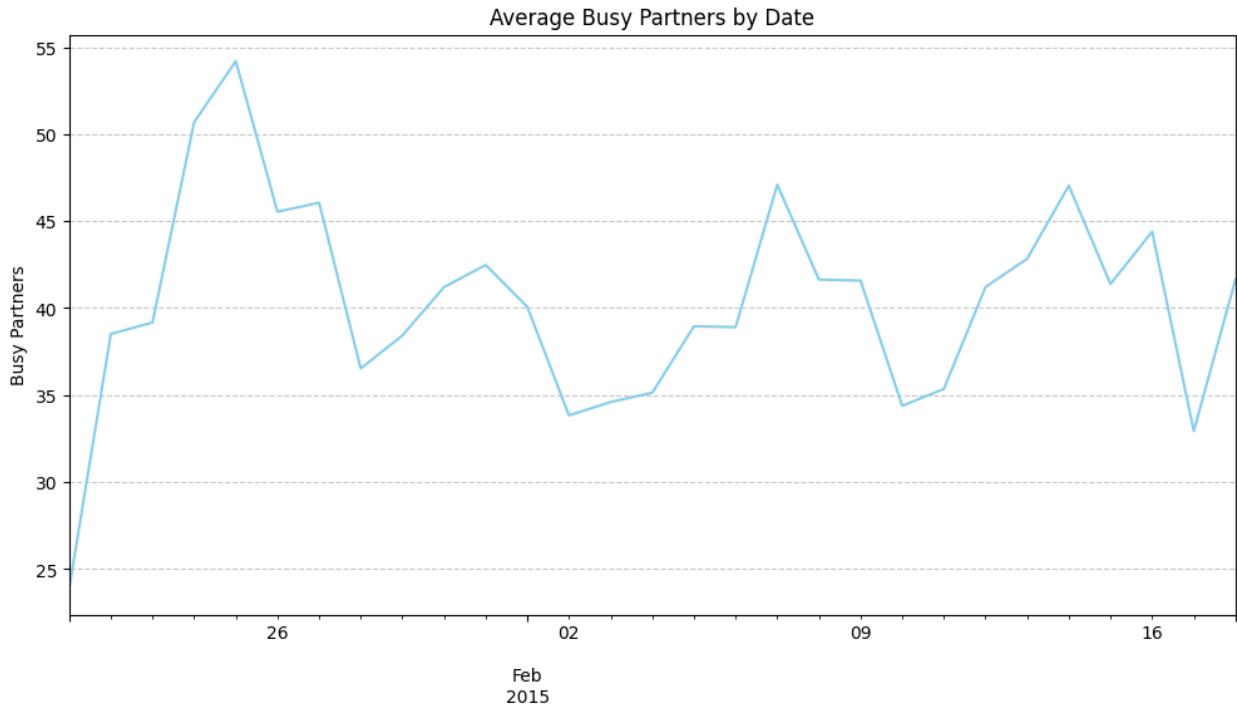


- Peak partner availability happens between 1 AM to 3 AM (66 at 2 AM), likely due to system or batch delivery times.
- Partner count drastically reduces post 4 AM, hitting lows around 7 AM to 8 AM.
- Evening hours (7 PM–9 PM) show a moderate rise again, reflecting increased demand during dinner-time deliveries.

Average Busy Partners by Date

```
# average busy partners by date
avg_busy_by_date = pdf.groupby('date')['total_busy_partners'].mean()

avg_busy_by_date.plot(kind='line', figsize=(12,6), title='Average Busy Partners by Date', color='skyblue')
plt.xlabel(None)
plt.ylabel('Busy Partners')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

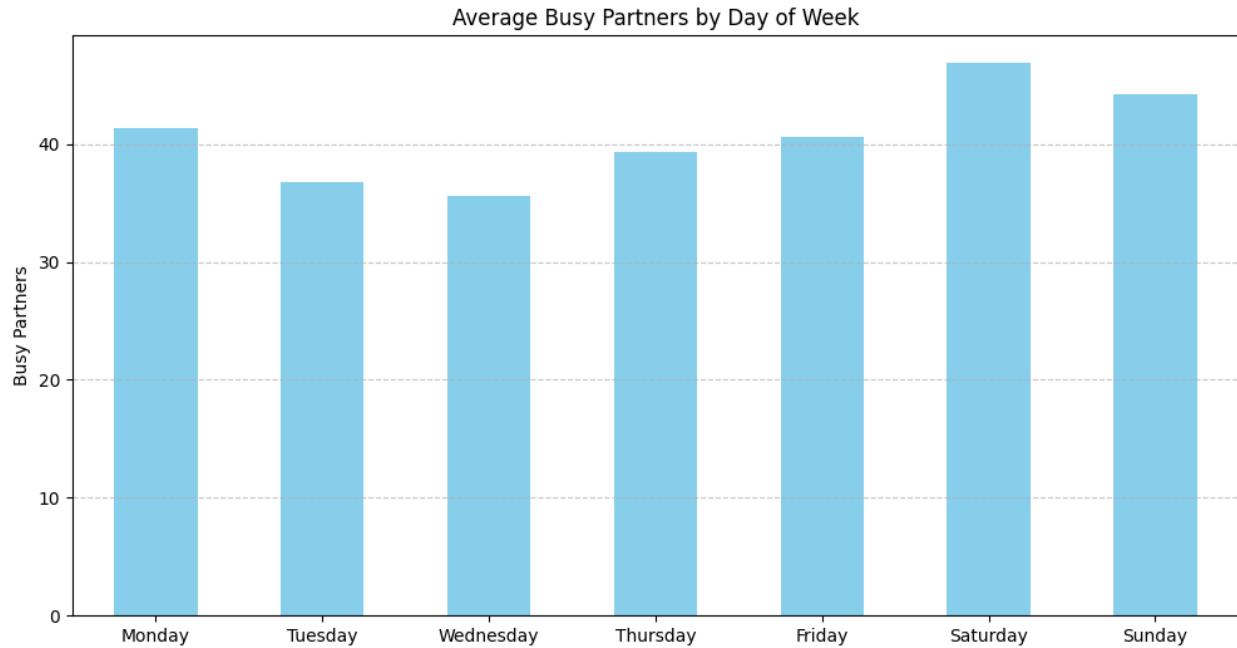


- The number of busy partners closely mirrors onshift trends, peaking around weekends.
- Highest busy partner count observed on Jan 25 (54) and Feb 7 (47).
- Some days like Feb 3–4 saw busy partner counts lower than onshift, hinting at potential overstaffing or lower demand.

Average Busy Partners by Day of Week

```
# average busy partners by day
avg_busy_by_day = (
    pdf.groupby('dayname')['total_busy_partners']
    .mean()
    .reindex([
        'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
        'Saturday', 'Sunday'
    ])
)

avg_busy_by_day.plot(kind='bar', figsize=(12, 6), title='Average Busy
Partners by Day of Week', color='skyblue')
plt.xlabel(None)
plt.ylabel('Busy Partners')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



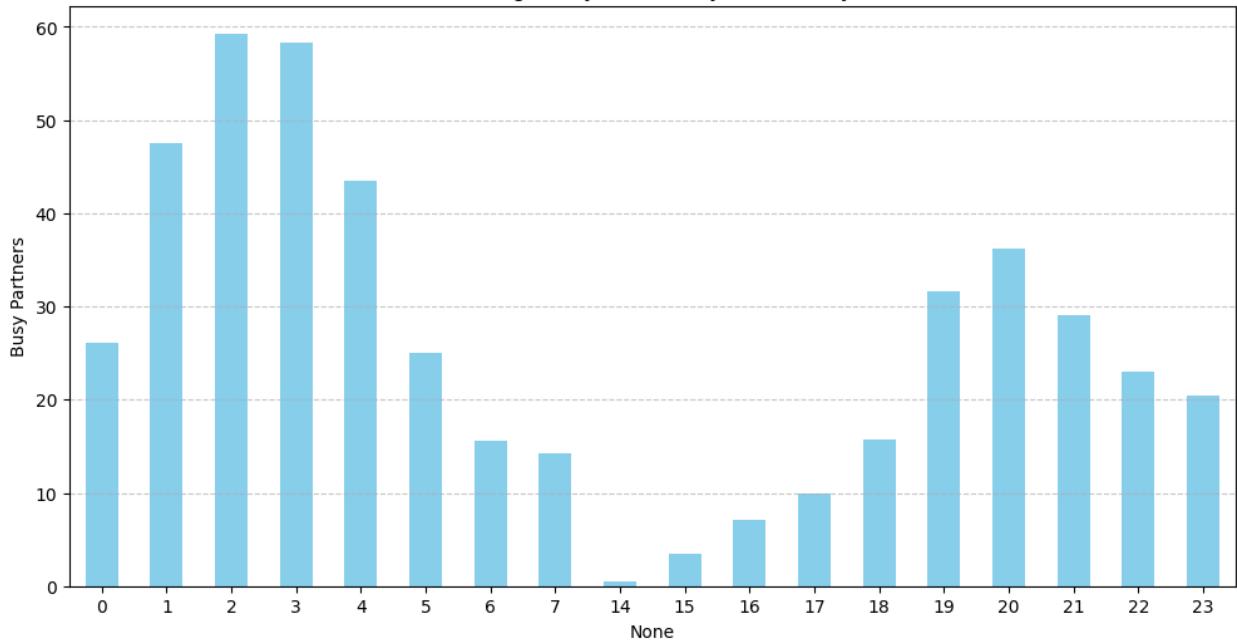
- Saturday (47) and Sunday (44) again stand out, confirming weekends as busiest.
- Wednesday (36) has the lowest average busy partners, aligned with lower delivery activity or supply-demand mismatch.
- The weekday dip suggests an opportunity to optimize resource allocation across the week.

Average Busy Partners by Hour of Day

```
# average busy partners by hour
avg_busy_by_hour = pdf.groupby('hour')[['total_busy_partners']].mean()

avg_busy_by_hour.plot(kind='bar', figsize=(12,6), title='Average Busy
Partners by Hour of Day', color='skyblue')
plt.xlabel('None')
plt.ylabel('Busy Partners')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=0)
plt.show()
```

Average Busy Partners by Hour of Day

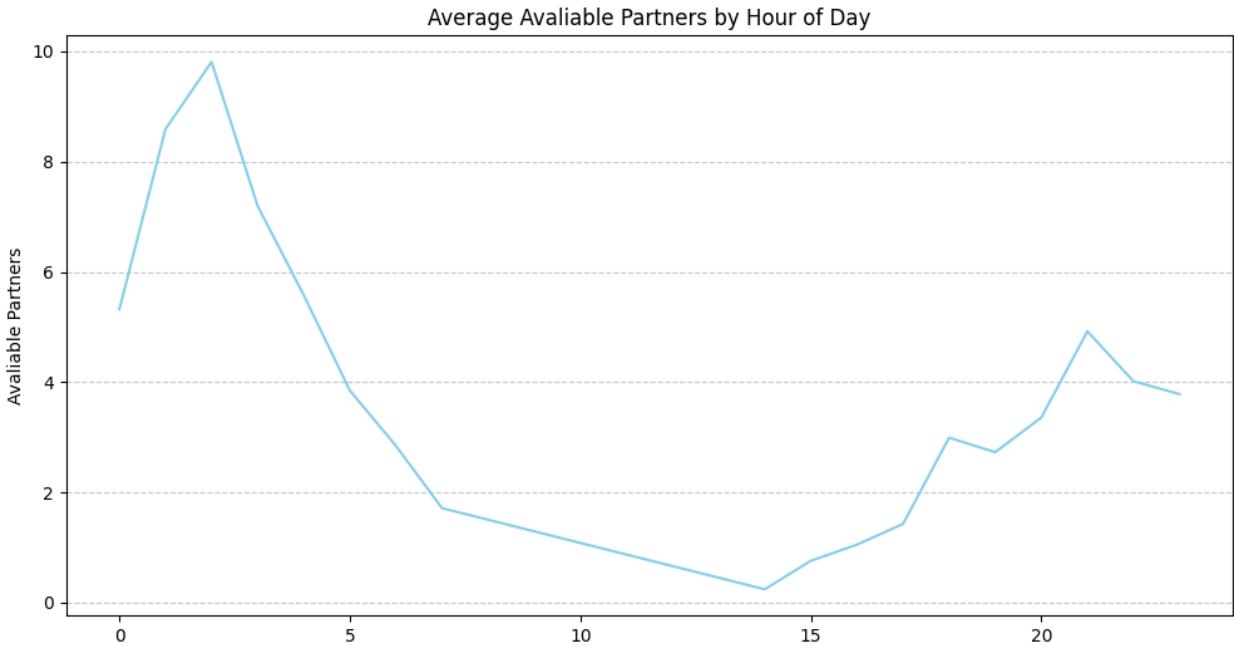


- Busy partner trends follow onshift trends closely, peaking at 2 AM (59) and staying high until 4 AM.
- Drop-off starts sharply post-4 AM, with the lowest engagement around 2 PM–5 PM.
- Evenings see a small resurgence, but not as high as early morning hours.

Average Available Partner by Hour of Day

```
# average available partners (onshift partners - busy partners) by hour
available_partner_by_hour = pdf.groupby('hour')[['available_partners']].mean()

available_partner_by_hour.plot(kind='line', figsize=(12,6),
                               title='Average Available Partners by Hour of Day', color='skyblue')
plt.xlabel(None)
plt.ylabel('Available Partners')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

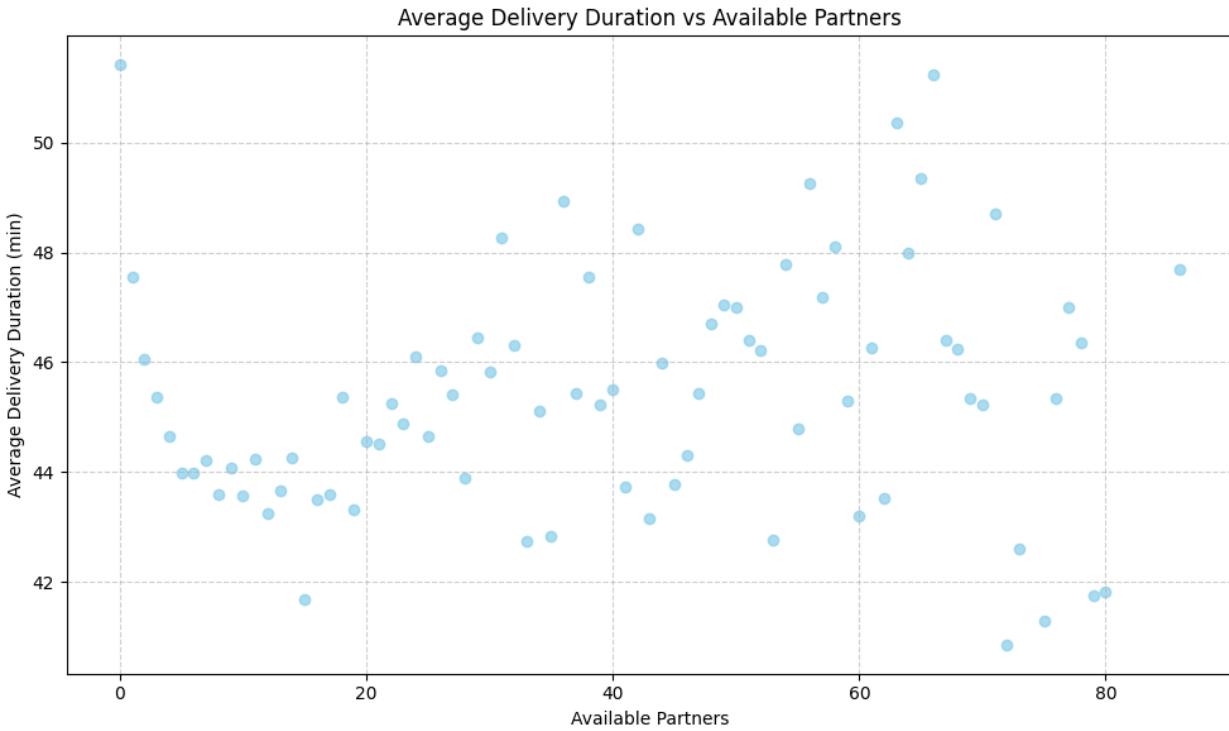


- Availability is highest between 2–3 AM, but even then, only about 9 partners are idle, showing high engagement levels.
- During evening hours (7–11 PM), available partners remain quite low (mostly 2–4), indicating a tightly stretched capacity.
- The lowest availability is seen during lunch-to-evening transition (2–6 PM), potentially a risk zone for delays.

Available Partners vs Average Delivery Duration

```
# available partners vs delivery duration
avg_duration_by_av_prt = pdf.groupby('available_partners')[['delivery_duration_min']].mean().reset_index()

plt.figure(figsize=(10, 6))
plt.scatter(avg_duration_by_av_prt['available_partners'],
            avg_duration_by_av_prt['delivery_duration_min'], color='skyblue',
            alpha=0.7)
plt.title('Average Delivery Duration vs Available Partners')
plt.xlabel('Available Partners')
plt.ylabel('Average Delivery Duration (min)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

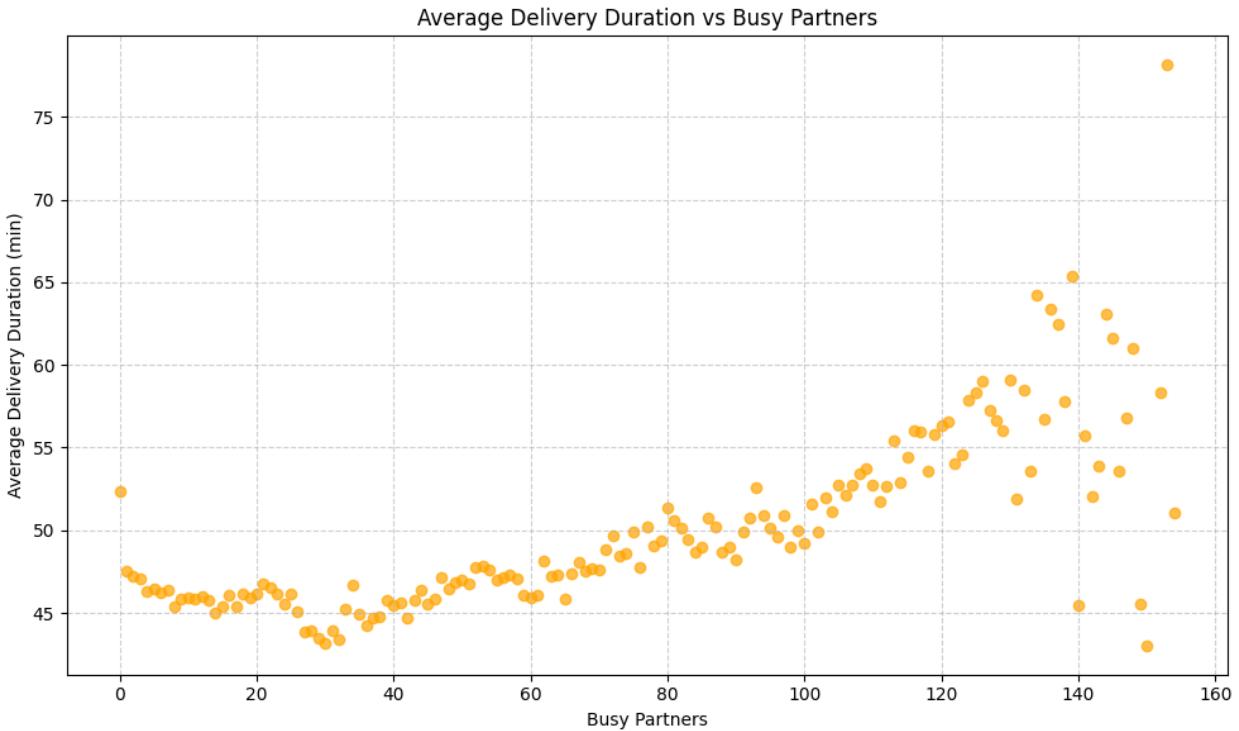


- When no partners are available, average delivery time is the highest at ~51.4 minutes, indicating possible delays due to resource unavailability.
- As the number of available partners increases, delivery duration gradually decreases, reaching around 41–45 minutes for optimal partner availability.
- Interestingly, even with 70+ partners available, delivery times plateau or slightly increase, implying diminishing returns or potential inefficiencies in partner utilization.
- This suggests a sweet spot exists where having more partners helps, but only up to a certain point – after which system coordination matters more.

Busy Partners vs Average Delivery Duration

```
# busy partners vs delivery duration
avg_duration_by_bu_prt = pdf.groupby('total_busy_partners')
['delivery_duration_min'].mean().reset_index()

plt.figure(figsize=(10, 6))
plt.scatter(avg_duration_by_bu_prt['total_busy_partners'],
avg_duration_by_bu_prt['delivery_duration_min'], color='orange',
alpha=0.7)
plt.title('Average Delivery Duration vs Busy Partners')
plt.xlabel('Busy Partners')
plt.ylabel('Average Delivery Duration (min)')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



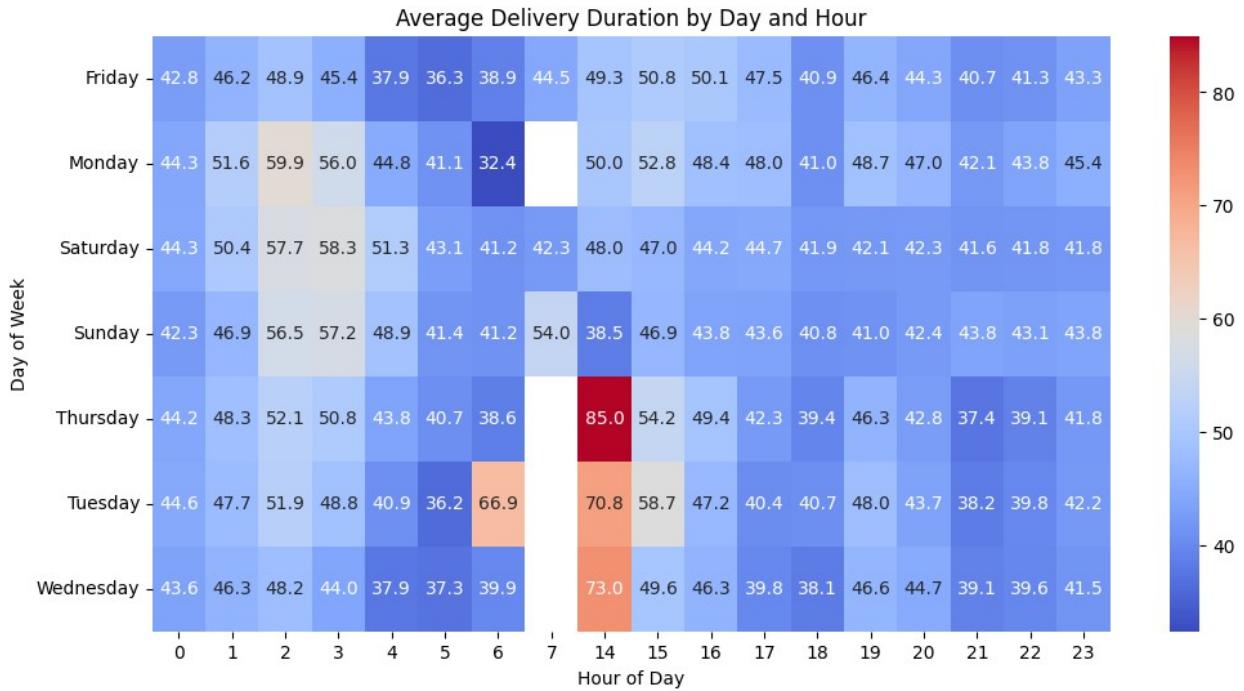
- Similar to the available partners trend, when zero busy partners are present, delivery duration is high (~52.3 minutes), likely due to lower operational activity.
- As more partners become busy (indicating more active deliveries), average delivery time decreases steadily to the mid-40-minute range.
- At very high busy partner counts (150+), delivery times become inconsistent, with a spike at 152 partners (58 mins) and 153 partners (78 mins), hinting at overload or system strain.
- This reflects that while more activity generally means better efficiency, extreme partner busyness may cause bottlenecks, impacting delivery performance.

Hour of Day vs Day of Week (Avg Delivery Duration & Order Volume)

```
# pivot table for average delivery duration
pivot_duration = pdf.pivot_table(values='delivery_duration_min',
index='dayname', columns='hour', aggfunc='mean')

plt.figure(figsize=(12, 6))

sns.heatmap(pivot_duration, cmap='coolwarm', annot=True, fmt=".1f")
plt.title('Average Delivery Duration by Day and Hour')
plt.ylabel('Day of Week')
plt.xlabel('Hour of Day')
plt.show()
```

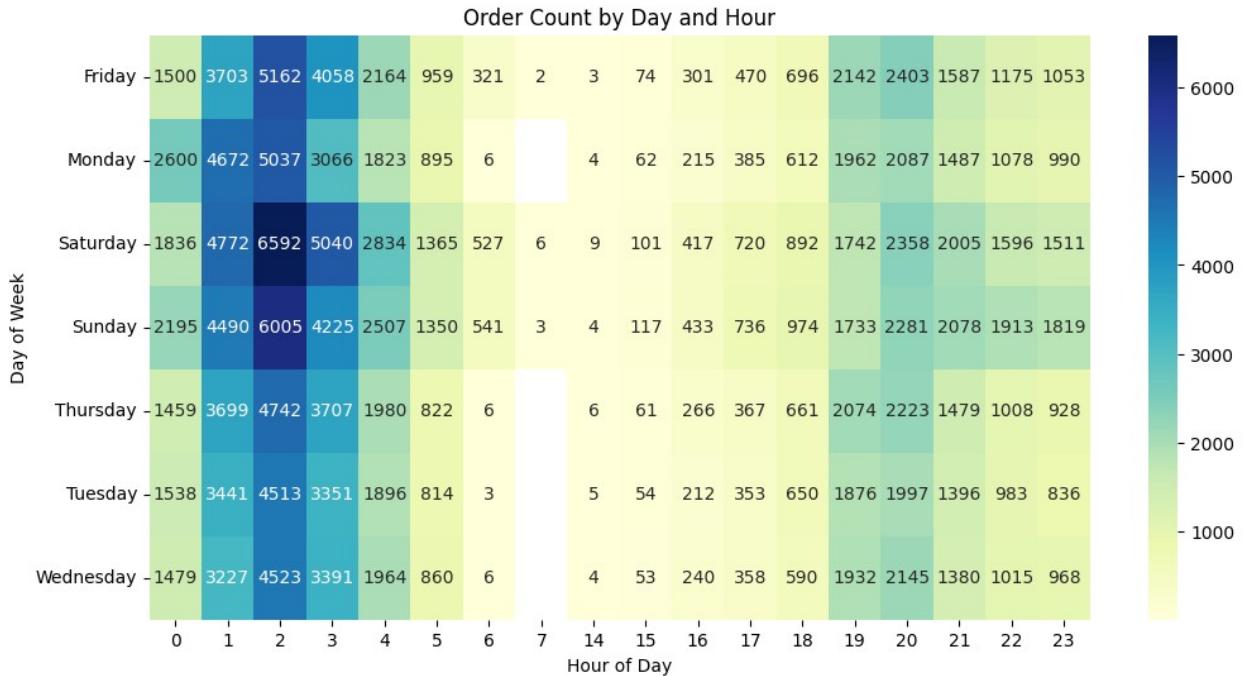


- Delivery durations are highest during these hours across all days, especially on weekends (Saturday and Sunday).
- This could be due to increased demand during late-night hours or reduced partner availability.
- Delivery durations are shortest during these hours on most days, indicating better resource allocation and lower demand variability.
- Monday and Saturday show consistently higher delivery durations across most hours, suggesting potential inefficiencies or higher demand on these days.
- Wednesday and Thursday have relatively lower delivery durations, indicating better operational efficiency midweek.
- Delivery durations spike significantly between 1 AM–4 AM, especially on weekends (Saturday and Sunday). This could be due to late-night orders or reduced partner availability.

```
# pivot table for order volume
pivot_orders = pdf.pivot_table(values='created_at', index='dayname',
columns='hour', aggfunc='count')

plt.figure(figsize=(12, 6))

sns.heatmap(pivot_orders, cmap='YlGnBu', annot=True, fmt=".0f")
plt.title('Order Count by Day and Hour')
plt.ylabel('Day of Week')
plt.xlabel('Hour of Day')
plt.show()
```



- Order volumes are highest during these hours, especially on weekends (Saturday and Sunday), reflecting late-night demand.
- Moderate order volumes are observed during these hours, likely due to dinner-time orders.
- Saturdays and Sundays generally have higher order volumes compared to weekdays, reflecting increased demand and potential strain on resources.
- Early morning hours (6 AM–2 PM) show the lowest order volumes, with a dip to almost zero between 7 AM–2 PM.

Available Partners vs Hour of Day with Delivery Duration

```
# available partners vs hour with duration
hourly = pdf.groupby('hour').agg({
    'available_partners': 'mean',
    'delivery_duration_min': 'mean'
}).reset_index()

fig, ax1 = plt.subplots(figsize=(10, 6))

color1 = 'tab:blue'
ax1.set_xlabel('Hour of Day')
ax1.set_ylabel('Available Partners', color=color1)
ax1.plot(hourly['hour'], hourly['available_partners'], color=color1,
label='Available Partners')
ax1.tick_params(axis='y', labelcolor=color1)

ax2 = ax1.twinx()
color2 = 'tab:red'
```

```

ax2.set_ylabel('Avg Delivery Duration (min)', color=color2)
ax2.plot(hourly['hour'], hourly['delivery_duration_min'],
color=color2, label='Avg Delivery Duration')
ax2.tick_params(axis='y', labelcolor=color2)

fig.suptitle('Available Partners vs Avg Delivery Duration by Hour')
fig.tight_layout()
plt.show()

```



- The late-night hours (1 AM–3 AM) show high demand and partner availability, but delivery durations remain relatively high, suggesting a need for better resource allocation or process optimization.
- Delivery durations are shortest during the evening hours (7 PM–9 PM), indicating efficient operations during this time.
- During periods of high partner availability (e.g., 2 AM), delivery durations do not decrease significantly, suggesting diminishing returns or inefficiencies in partner utilization.

Recommendations

Dynamic Partner Allocation

- Late-night (1–4 AM) and weekend demand surges cause delays due to partner shortages.
- When more delivery partners are available, deliveries are faster.
- If there are very few or no partners available, delivery times increase.
- Move delivery partners to areas where more orders are expected, especially during midnight and dinner.

- In short, boost partner availability during weekends (Fri–Sun) and late nights (1–4 AM) with: Bonus pay for delivery partners, using surge pricing for customers (to manage demand)

Restaurant Level Interventions

- Some food types get more orders, like American food, Pizza, Mexican.
- Since these food types are in high demand, they may take longer to prepare.
- Slow prep times for certain cuisines (Caribbean, Belgian) and bulk orders (36+ items).
- Suggested pre-prep ingredients during peak hours.
- Add a fee for orders >10 items to compensate for extra delivery time.

Order Protocol Optimization

- Protocol 6 has 59-min avg. delivery time vs. 41 mins for Protocol 7.
- Investigate if it's a legacy system (e.g., call-in orders) and migrate users to faster channels (app/web).
- Offer discounts for customers using Protocol 3 (avg. 45 mins).

Demand Shaping & Customer Communication

- 2–4 AM orders have high demand but long delivery times (54+ mins).
- Add a 10% fee for 1–4 AM orders to smooth demand.
- Incentivize 6–10 AM orders with 15% off to balance demand.

Operational Efficiency Fixes

- Mondays/Saturdays have 50+ min deliveries despite moderate demand.
- Audit traffic routes on these days for bottlenecks.
- Assign dedicated "fast-response" partners for high-priority orders.

Data-Driven Partner Thresholds

- Delivery times spike when available partners <5 or outstanding orders >50.
- Maintain 10% idle partners during peak hours for flexibility.